

# REVE1

## Translation of C to Assembly

```
absdiff:
    cmpq    %rsi, %rdi
    jg      .L4
    movq    %rsi, %rax
    subq    %rdi, %rax
    jmp     .L7
.L4:
    movq    %rdi, %rax
    subq    %rsi, %rax
.L7:
    ret
```

# Module Outline

## ■ Assembly Programming

- Program Structure
- Language Elements

## ■ Altering Control Flow

- If-then-else Constructs
- Loop Constructs
- Switch Statements

## ■ Composite Data Structures

- Arrays
- Structures
- Unions

## ■ Calling Convention

## ■ Module Summary

# Assembly Programming

## Program Structure

```
.text
.align 16

.globl main

main:
    movl $1, %eax
    movl $1, %edi
    leaq hstr(%rip), %rsi
    movl $14, %edx

    syscall

    xor    %eax, %eax
    ret

.data
.align 16

hstr:
    .ascii "Hello, world!\n"
```

# Hello, world!

## ■ Our first assembly program

```
.text
.align 16

.globl main

main:
    movl    $1, %eax
    movl    $1, %edi
    leaq    hstr(%rip), %rsi
    movl    $14, %edx

    syscall

    xor     %eax, %eax
    ret

.data
.align 16

hstr:
    .ascii "Hello, world!\n"
```

# Hello, world!

## ■ Our first assembly program

```
$ gcc -o hello hello.S  
$ ./hello  
Hello, world!  
$
```

# General Structure

<code>.text</code>	{ section specifier (.text = code)
<code>.align 16</code>	{ alignment specification
<code>.globl main</code>	{ exported symbols
<code>main:</code> <code>movl \$1, %eax</code> <code>...</code> <code>ret</code>	{ function definitions
<code>.data</code>	{ section specifier (.data = global variables)
<code>.align 16</code>	{ alignment specification
<code>hstr:</code>	{ label (memory address alias, i.e., variable)
<code>.ascii "Hello, world!\n"</code>	{ data (value of initialized variable)

# Assembly Programming

## Language Elements

# The GNU Assembler (GAS)

## ■ The GNU Assembler

- gas, as
- assembler maintained by the [GNU project](#)
- supports many, many different architectures

## ■ Syntax varies by architecture

- to match SOP of those architectures
- last line must be terminated by a newline

## ■ Many, many options and directives

- refer to the excellent [GAS manual](#)



# Language Elements

## ■ Block delimiter

- no blocks in assembly

## ■ Whitespace

- spaces, tabs, newlines
- is ignored
- no forced indentation

## ■ Command end marker

- newline
- last line terminated by a newline

## ■ Comments

- single line: #
- multi-line: /\* ... \*/
- multi-line > single line
- nesting of same-level comments not supported

```
.text
.align 16

.globl main

/*
 * main function
 */
main:
    movl  $1, %eax    # set %rax to 1
    ...
    ret               # return to caller

.data
.align 16

hstr:
    .ascii "Hello, world!\n"
```

# Language Elements

## ■ Symbols

- used to name things
  - ▶ code locations
    - functions
    - jump targets
  - ▶ data
    - variables
- the dot symbol
  - ▶ refers to the current location
  - ▶ very handy to compute static string length
- syntax

```
.text
.align 16

.globl main

/*
 * main function
 */
main:
    movl    $1, %eax    # set %rax to 1
    ...
    leaq    hstr(%rip), %rsi
    movl    $hlen, %edx
    ...
    ret                                # return to caller

.data
.align 16

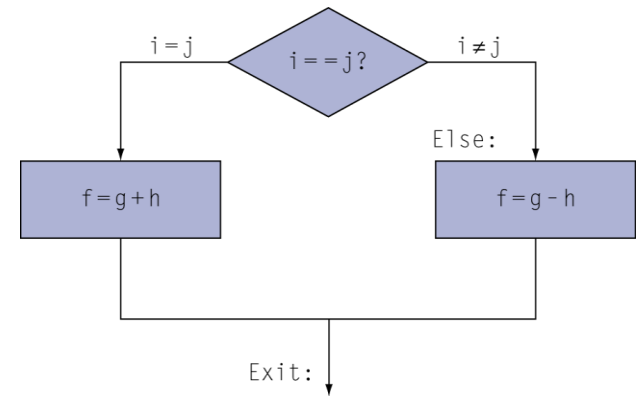
hstr:
    .ascii "Hello, world!\n"
    hlen = . - hstr
```

# Language Elements

## ■ Directives

- large number of directives
- differ by architecture
- **most commonly used directives**

.text	.ascii
.data	.asciz
.align	.byte
.extern	.int
.globl	.long
	.quad
	.short
	.size
	.skip
	.string
	.word
	.zero



# Altering Control Flow

# Altering Control Flow

- Higher-level programming languages offer control-flow altering constructs

```
int foo(int x, int y)
{
    int res = 0;

    if (x > y) x = x-y;

    while (x > 0) {
        res = res + y;
        x--;
    }

    return res;
}
```

- How can we achieve that with assembly code?

# Branch Operations

## ■ Processor ISAs offer branch operations to alter the sequential control flow

- generic form

```
branch <label>
```

- instructs processor to continue execution at <label>
  - ▶ same as goto in higher-level programming languages

- branch <label> implemented as  
PC = &label

- **unconditional branch**

- ▶ branch is always executed

- differences by architectures

- ▶ Intel:

```
jmp <label>
```

- ▶ RISC-V:

```
branch <label>
```

```
...  
subl    %eax, %edx  
movl    %edx, %eax  
jmp     .L7  
...
```

```
.L7:
```

```
...
```

# Branch Operations

## ■ Conditional branches

- *conditionally* alter control flow
- generic form

```
branch <condition>, <label>
```

- instructs processor to continue execution at <label> if <condition> is true
  - ▶ if (condition) goto label

- different ways to implement conditional branches

```
if (operand1 <cond> operand2) goto <label>
```

- ▶ Intel:

```
cmp    operand2, operand1  
j<cond> <label>
```

- ▶ RISC-V:

```
b<cond> operand1, operand2, <label>
```

# Recap: x86 Condition Codes

## ■ Single bit registers

CF	Carry Flag (for unsigned)
SF	Sign Flag (for signed)
ZF	Zero Flag
OF	Overflow Flag (for signed)

## ■ Implicitly set by arithmetic operations

Example: `addq/addl src,dest`  $\Leftrightarrow$  `t = a+b`

**CF set** if carry out from most significant bit (unsigned overflow)

**ZF set** if `t == 0`

**SF set** if `t < 0` (as signed)

**OF set** if two's-complement (signed) overflow  
(`a>0 && b>0 && t<0`) || (`a<0 && b<0 && t>=0`)

- Not set by `lea` instruction

## ■ Explicitly set by

- `cmp S2, S1` set condition codes based on **S1-S2**
- `test S2, S1` set condition codes based on **S1 & S2**

%rax	%eax
%rbx	%ebx
%rcx	%ecx
%rdx	%edx
%rsi	%esi
%rdi	%edi
%rsp	%esp
%rbp	%ebp
%r8	%r8d
%r9	%r9d
%r10	%r10d
%r11	%r11d
%r12	%r12d
%r13	%r13d
%r14	%r14d
%r15	%r15d

%rip

Instruction pointer  
(read-only access)

CF	ZF	SF	OF
----	----	----	----

Condition codes



# Recap: x86 Condition Codes/Jumps

- setX: reading condition codes
  - set single byte, does not alter remaining 3 bytes
- jX: jump to different part of code depending on condition codes
- cmovX: conditional move

setX dest	jX target	cmovX S, reg	Condition	Description
	jmp		1	
sete	je	cmove	ZF	Equal / Zero
setne	jne	cmovne	$\sim$ ZF	Not Equal / Not Zero
sets	js	cmovs	SF	Negative
setns	jns	cmovns	$\sim$ SF	Nonnegative
setg	jg	cmovg	$\sim$ (SF $\wedge$ OF)& $\sim$ ZF	Greater (Signed)
setge	jge	cmovge	$\sim$ (SF $\wedge$ OF)	Greater or Equal (Signed)
setl	jl	cmovl	(SF $\wedge$ OF)	Less (Signed)
setle	jle	cmovle	(SF $\wedge$ OF)   ZF	Less or Equal (Signed)
seta	ja	cmova	$\sim$ CF& $\sim$ ZF	Above (unsigned)
setb	jb	cmovb	CF	Below (unsigned)

```
...  
movq    %rsi, %rdi  
cmpq    %rsi, %r8  
jl      .L2  
movq    %r8, %rdi  
call    foo@PLT  
jmp     .L3  
...
```

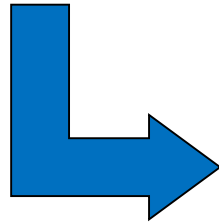
## Altering Control Flow

# If-then-else Constructs

# Compiling If Statements

**C code:**

```
int cf_if(long a, long b) {  
    int res;  
  
    if (a >= b) res = foo(a, b);  
    else res = bar(b, a);  
  
    return res*2;  
}
```



**C code (goto version):**

```
    if (a < b) goto L_false;  
    res = foo(a, b)  
    goto Exit;  
  
L_false:  
    res = bar(b, a)  
  
Exit:  
    return res*2;
```

# Compiling If Statements

## C code:

```
int cf_if(long a, long b) {
    int res;

    if (a >= b) res = foo(a, b);
    else res = bar(b, a);

    return res*2;
}
```

## C code (goto version):

```
if (a < b) goto L_false;
res = foo(a, b)
goto Exit;

L_false:
    res = bar(b, a)

Exit:
    return res*2;
```

## Direct translation to x86 code:

```
cf_if:
    movq    %rdi, %r8
    subq    $8, %rsp
    movq    %rsi, %rdi
    cmpq    %rsi, %r8
    jl      .L_false
    movq    %r8, %rdi
    call    foo@PLT
    jmp     .L_exit

.L_false:
    movq    %r8, %rsi
    call    bar@PLT

.L_exit:
    addq    $8, %rsp
    addl    %eax, %eax
    ret
```

# Compiling If Statements

## C code:

```
int cf_if(long a, long b) {
    int res;

    if (a >= b) res = foo(a, b);
    else res = bar(b, a);

    return res*2;
}
```

## C code (goto version):

```
if (a < b) goto L_false;
res = foo(a, b)
goto Exit;

L_false:
    res = bar(b, a)

Exit:
    return res*2;
```

## Direct translation to x86 code:

```
cf_if:
    movq    %rdi, %r8
    subq    $8, %rsp
    movq    %rsi, %rdi
    cmpq    %rsi, %r8
    jl      .L_false
    movq    %r8, %rdi
    call    foo@PLT
    jmp     .L_exit

.L_false:
    movq    %r8, %rsi
    call    bar@PLT

.L_exit:
    addq    $8, %rsp
    addl    %eax, %eax
    ret
```

# Compiling If Statements

## Code generated by GCC 9.3

```
cf_if:
    movq    %rdi, %r8
    subq    $8, %rsp
    movq    %rsi, %rdi
    cmpq    %rsi, %r8
    jl      .L2
    movq    %r8, %rdi
    call    foo@PLT
    addq    $8, %rsp
    addl    %eax, %eax
    ret
.L2:
    movq    %r8, %rsi
    call    bar@PLT
    addq    $8, %rsp
    addl    %eax, %eax
    ret
```

## ■ Observations

- minimize control flow instructions
- minimize instructions in general
- aligns stack points at 16-byte boundaries
- condition sometimes reversed
  - ▶ and fix jumps to if/else bodies

```
for (Init; Test; Update)  
  Body
```



```
Init;  
while (Test) {  
  Body  
  Update;  
}
```



```
Init;  
if (!Test)  
  goto done;  
do  
  Body  
  Update  
while (Test);  
done:
```



```
Init;  
if (!Test)  
  goto done;  
loop:  
  Body  
  Update  
  if (Test)  
    goto loop;  
done:
```

# Altering Control Flow

## Loop Constructs

# Loop Statements

## ■ Basic loop constructs

- `do { body } while (cond);`
- `while (cond) { body }`
- `for (init; cond; update) { body }`

```
do {  
    i += 1;  
} while (A[i] < i);
```

```
while (A[i] == k) {  
    i += 1;  
};
```

```
for (i=0; i<N; i++) {  
    sum += A[i];  
}
```



# General “Do-While” Translation

## C Code

```
do
    Body
while (Test);
```

## Goto Version

```
loop:
    Body
    if (Test)
        goto loop;
```

- Body: 

```
{
    Statement1;
    Statement2;
    ...
    Statementn;
}
```
- Test returns integer
  - = 0 interpreted as false
  - ≠ 0 interpreted as true

# General “While” Translation

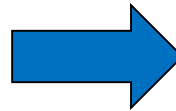
While version

```
while (Test)  
    Body
```



Do-While Version

```
if (!Test)  
    goto done;  
do  
    Body  
while (Test);  
done:
```



Goto Version

```
if (!Test)  
    goto done;  
loop:  
    Body  
    if (Test)  
        goto loop;  
done:
```

# Loop Statements

## ■ While as Do-While

- initial check
- move condition to end

```
while (A[i] == k) {  
    i += 1;  
};
```



```
if (A[i] != k) goto Exit;  
do {  
    i += 1;  
} while (A[i] == k);  
Exit:
```

# “For” Loop → While → Do While → Goto

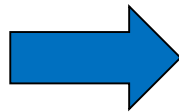
## For Version

```
for (Init; Test; Update)  
  Body
```



## While Version

```
Init;  
while (Test) {  
  Body  
  Update;  
}
```

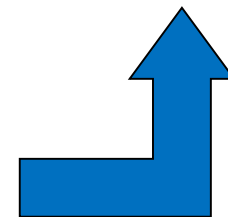


## Do-While Version

```
Init;  
if (!Test)  
  goto done;  
do  
  Body  
  Update  
while (Test);  
done:
```

## Goto Version

```
Init;  
if (!Test)  
  goto done;  
loop:  
  Body  
  Update  
  if (Test)  
    goto loop;  
done:
```



# “For” Loop Conversion Example

## C Code

```
#define WSIZE 8*sizeof(int)
int pcount_for(unsigned x) {
    int i;
    int result = 0;
    for (i = 0; i < WSIZE; i++) {
        unsigned mask = 1 << i;
        result += (x & mask) != 0;
    }
    return result;
}
```

- In this case, the initial test can be optimized away

## Goto Version

```
int pcount_for_gt(unsigned x) {
    int i;
    int result = 0;
    i = 0;
    if (!(i < WSIZE))
    goto done;
loop:
    {
        unsigned mask = 1 << i;
        result += (x & mask) != 0;
    }
    i++;
    if (i < WSIZE)
        goto loop;
done:
    return result;
}
```

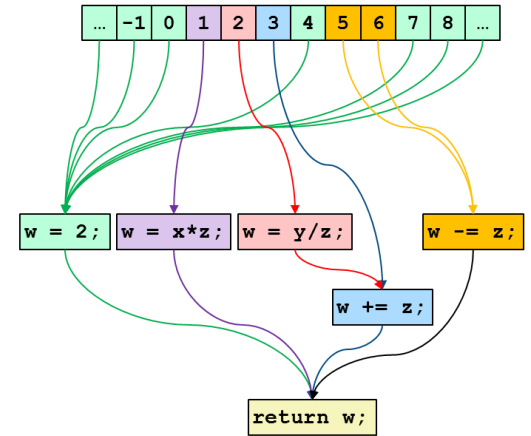
*Init*

*! Test*

*Body*

*Update*

*Test*



# Altering Control Flow

## Switch Statements

```

long switch_eg
(long x, long y, long z)
{
    long w = 1;
    switch(x) {
        case 1:
            w = y*z;
            break;
        case 2:
            w = y/z;
            // fall through
        case 3:
            w += z;
            break;
        case 5:
        case 6:
            w -= z;
            break;
        default:
            w = 2;
    }
    return w;
}

```

## Switch Statement Example

- Multiple case labels (here: 5 & 6)
- Fall through cases (here: 2)
- Missing cases (here: 4)

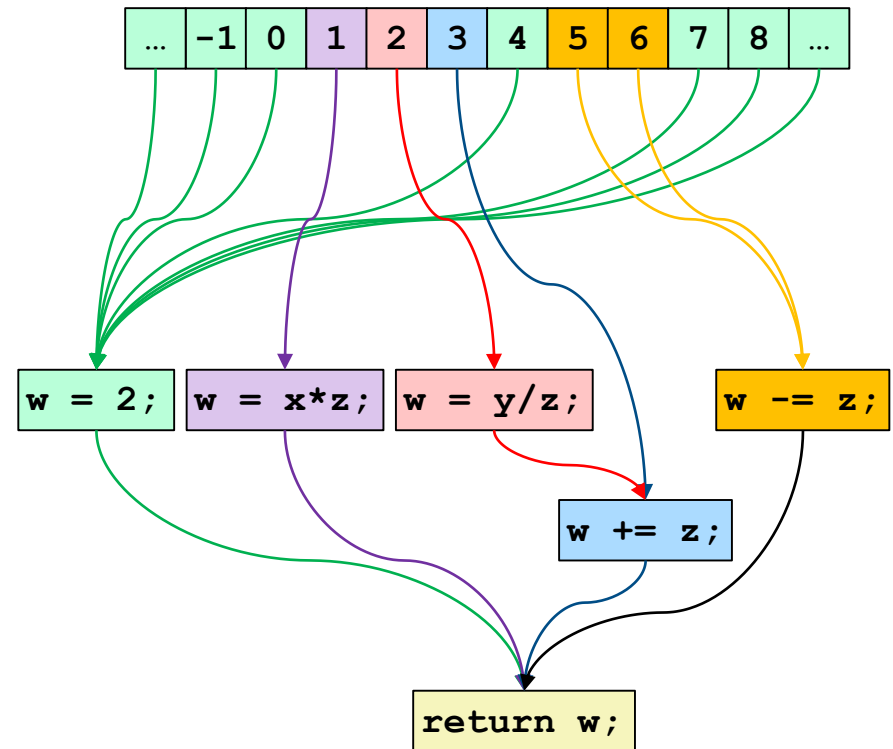
```

long switch_eg
(long x, long y, long z)
{
    long w = 1;
    switch(x) {
        case 1:
            w = y*z;
            break;
        case 2:
            w = y/z;
            // fall through
        case 3:
            w += z;
            break;
        case 5:
        case 6:
            w -= z;
            break;
        default:
            w = 2;
    }
    return w;
}

```

# Switch Statement Example

- Execution flow in dependence of **x**





```

long switch_eg
(long x, long y, long z)
{
    long w = 1;
    switch(x) {
        case 1:
            w = y*z;
            break;
        case 2:
            w = y/z;
            // fall through
        case 3:
            w += z;
            break;
        case 5:
        case 6:
            w -= z;
            break;
        default:
            w = 2;
    }
    return w;
}

```

## (Inefficient) Implementation using nested if statements

```

long switch_eg
(long x, long y, long z)
{
    long w = 1;
    if (x == 1) {
        w = y*z;
    } else {
        if (x == 2) {
            w = y/z;
        } else {
            if (x == 3) {
                w += z;
            } else {
                if ((x == 5) || (x == 6)) {
                    w -= z;
                } else {
                    w = 2;
                }
            }
        }
    }
    return w;
}

```

```

long switch_eg
(long x, long y, long z)
{
    long w = 1;
    switch(x) {
        case 1:
            w = y*z;
            break;
        case 2:
            w = y/z;
            // fall through
        case 3:
            w += z;
            break;
        case 5:
        case 6:
            w -= z;
            break;
        default:
            w = 2;
    }
    return w;
}

```

## Efficient implementation using a jump table

```

long switch_eg
(long x, long y, long z)
{
    long w = 1;

    if ((x < 1) || (x > 6)) goto default;
    goto jtab[x];

case1:
    w = y*z;
    goto end;
case2:
    w = y/z;
case3:
    w += z;
    goto end;
case56:
    w -= z;
    goto end;
default:
    w = 2;
end:
    return w;
}

```

jtab	
index	target
0	default
1	case1
2	case2
3	case3
4	default
5	case56
6	case56

# Efficient Implementation using a Jump Table

## Switch Form

```
switch(x) {  
  case val_0:  
    Block 0  
  case val_1:  
    Block 1  
    . . .  
  case val_n-1:  
    Block n-1  
}
```

## Jump Table

jtab:	Targ0
	Targ1
	Targ2
	•
	•
	•
	Targn-1

## Jump Targets

Targ0:

Code Block  
0

Targ1:

Code Block  
1

Targ2:

Code Block  
2

•  
•  
•

Targn-1:

Code Block  
n-1

## Approximate Translation

```
target = jtab[x];  
goto *target;
```

# Switch Statement Example (x86\_64)

```
long switch_eg(long x, long y, long z)
{
    long w = 1;
    switch(x) {
        . . .
    }
    return w;
}
```

## Setup:

switch\_eg:

```
movq    %rdx, %rcx
cmpq    $6, %rdi
ja      .L8
```

```
# rcx = copy of z
# compare x and 6
# if above (unsigned >) goto L8
```



What range of values  
takes default?

# Switch Statement Example (x86\_64)

```
long switch_eg(long x, long y, long z)
{
    long w = 1;
    switch(x) {
        . . .
    }
    return w;
}
```

## Jump table

.section	.rodata
.align 8	
.L4:	
.quad	.L8 # x = 0
.quad	.L3 # x = 1
.quad	.L5 # x = 2
.quad	.L9 # x = 3
.quad	.L8 # x = 4
.quad	.L7 # x = 5
.quad	.L7 # x = 6

## Setup:

```
switch_eg:
    movq    %rdx, %rcx    # rcx = copy of z
    cmpq    $6, %rdi      # compare x and 6
    ja      .L8           # if above (unsigned >) goto L8
    jmp     *.L4(,%rdi,8)  # goto *jtab[x]
```

- Notes:
- (1) `w` not initialized
  - (2) `jmp *<address>` executes an indirect jump

# Assembly Setup Explanation

## ■ Table Structure

- Each target requires 8 bytes
- Base address at .L4

## ■ Jumping

- **Direct:** `jmp .L2`
- Jump target is denoted by label .L2
- **Indirect:** `jmp *.L4(,%rdi,8)`
- Start of jump table: .L4
- Must scale by factor of 8 (labels have 64-bits = 8 bytes addresses on x86\_64)
- Fetch target from effective Address `.L4 + rdi*8`
  - ▶ Only for  $0 \leq x \leq 6$

## Jump table

```
.section      .rodata
    .align 8
.L4:
    .quad     .L8 # x = 0
    .quad     .L3 # x = 1
    .quad     .L5 # x = 2
    .quad     .L9 # x = 3
    .quad     .L8 # x = 4
    .quad     .L7 # x = 5
    .quad     .L7 # x = 6
```

# Jump Table

## Jump table

```
.section .rodata
.align 8
.L4:
.quad .L8 # x = 0
.quad .L3 # x = 1
.quad .L5 # x = 2
.quad .L9 # x = 3
.quad .L8 # x = 4
.quad .L7 # x = 5
.quad .L7 # x = 6
```

```
switch(x) {
case 1: // .L3
    w = y*z;
    break;
case 2: // .L5
    w = y/z;
    // fall through
case 3: // .L9
    w += z;
    break;
case 5: case 6: // .L7
    w -= z;
    break;
default: // .L8
    w = 2;
}
```

# Handling Fall-Through

```
long w = 1;  
...  
switch(x) {  
  ...  
  case 2:   
    w = y/z;  
    // fall through  
  case 3:   
    w += z;  
    break;  
  ...  
}
```

case 3:  
w = 1;  
goto merge;

case 2:  
w = y/z;  
merge:  
w += z;



# Full Example

```
long switch_eg
(long x, long y, long z)
{
    long w = 1;
    switch(x) {
        case 1:
            w = y*z;
            break;
        case 2:
            w = y/z;
            // fall through
        case 3:
            w += z;
            break;
        case 5:
        case 6:
            w -= z;
            break;
        default:
            w = 2;
    }
    return w;
}
```

```
switch_eg:
    movq    %rdx, %rcx
    cmpq    $6, %rdi
    ja      .L8
    jmp     *.L4(,%rdi,8)
.L3:
    movq    %rsi, %rax
    imulq   %rdx, %rax
    ret
.L5:
    movq    %rsi, %rax
    cqto
    idivq   %rcx
    jmp     .L6
.L9:
    movl    $1, %eax
.L6:
    addq    %rcx, %rax
    ret
.L7:
    movl    $1, %eax
    subq    %rdx, %rax
    ret
.L8:
    movl    $2, %eax
    ret
```

```
.L4:
    .quad   .L8
    .quad   .L3
    .quad   .L5
    .quad   .L9
    .quad   .L8
    .quad   .L7
    .quad   .L7
```

# Understanding the Assembly Code

```
long switch_eg
(long x, long y, long z)
{
    rdi    rsi    rdx
    long w = 1;
    switch(x) {
        case 1:
            w = y*z;
            break;
        case 2:
            w = y/z;
            /* Fall Through */
        case 3:
            w += z;
            break;
        case 5:
        case 6:
            w -= z;
            break;
        default:
            w = 2;
    }
    return w;
}
```

```
switch_eg:
    movq    %rdx, %rcx
    cmpq    $6, %rdi
    ja      .L8
    jmp     *.L4(, %rdi, 8)
.L3:
    movq    %rsi, %rax
    imulq   %rdx, %rax
    ret
.L5:
    movq    %rsi, %rax
    cqto
    idivq   %rcx
    jmp     .L6
.L9:
    movl    $1, %eax
.L6:
    addq    %rcx, %rax
    ret
.L7:
    movl    $1, %eax
    subq    %rdx, %rax
    ret
.L8:
    movl    $2, %eax
    ret
```

```
.L4:
    .quad   .L8 # x = 0
    .quad   .L3 # x = 1
    .quad   .L5 # x = 2
    .quad   .L9 # x = 3
    .quad   .L8 # x = 4
    .quad   .L7 # x = 5
    .quad   .L7 # x = 6
```

# Understanding the Assembly Code

■ switch.s: complete output of `gcc -m64 -O -S switch.c`

- jump table embedded in code
- separated into code and data section by assembler

```
.file      "switch.c"
.text
.globl    switch_eg
.type     switch_eg, @function
switch_eg:
.LFB0:
.cfi_startproc
movq      %rdx, %rcx
cmpq      $6, %rdi
ja        .L8
jmp        *.L4(,%rdi,8)
.section   .rodata
.align    8
.align    4
.L4:
.quad     .L8
.quad     .L3
.quad     .L5
.quad     .L9
.quad     .L8
.quad     .L7
.quad     .L7
.text
.L3:
movq      %rsi, %rax
imulq     %rdx, %rax
ret
```

```
.L5:
movq      %rsi, %rax
cqto
idivq     %rcx
jmp        .L6
.L9:
movl      $1, %eax
.L6:
addq      %rcx, %rax
ret
.L7:
movl      $1, %eax
subq      %rdx, %rax
ret
.L8:
movl      $2, %eax
ret
.cfi_endproc
.LFE0:
.size     switch_eg, .
.ident    "GCC: (Gentoo
4.9.4 p1.0,
pie-0.6.4) 4.9.4"
.section   .note.GNU-stack,
""@progbits
```

# Understanding the Assembly Code

## ■ Noteworthy Features

- Jump table avoids sequencing through cases
  - ▶ constant time, rather than linear
- Use jump table to handle holes and duplicate tags
- Use program sequencing to handle fall-through
- Compiler decides low-level implementation of switch (jump table, if-elsif-elsif-else, decision tree)

# Object Code

## Assembly Code

```
switch_eg:
...
ja    .L8          # if above (unsigned >) goto default
jmp   *.L4(,%rdi,8) # goto *jtab[x]
```

## Disassembled Object Code

```
0000000000400508 <switch_eg>:
...
40050f: 77 07          ja    40053c <switch_eg+0x34>
400511: ff 24 fd ...   jmp   *0x4005c8(,%rdi,8)
```

- Matching assembly to object code
  - code of default (**.L8**) located at address **0x000000000040053c**
  - jump table (**.L4**) located at address **0x00000000004005c8**

# Object Code: Jump Table

- jump table does not show up in disassembled code
  - located in *data* section
  - can use `objdump -dD <file>` to also disassemble data sections, but result is not very useful

- better using GDB (GNU debugger)

```
$ gdb switch
```

```
(gdb) x/7xg 0x4005c8
```

```
0x4005c8:      0x00000000000040053c      0x000000000000400518
```

```
0x4005d8:      0x000000000000400520      0x00000000000040052a
```

```
0x4005e8:      0x00000000000040053c      0x000000000000400533
```

```
0x4005f8:      0x000000000000400533
```

```
(gdb)
```

- `x/7xg`: examine 7 hexadecimal format “giant words” (8 bytes)
- Use command “`help x`” to get format documentation

# Object Code: Deciphering the Jump Table

```
(gdb) x/7xg 0x4005c8
```

```
0x4005c8:    0x000000000040053c    0x0000000000400518
0x4005d8:    0x0000000000400520    0x000000000040052a
0x4005e8:    0x000000000040053c    0x0000000000400533
0x4005f8:    0x0000000000400533
```

Address	Value	x
0x4005c8	0x40053c	0
0x4005d0	0x400518	1
0x4005d8	0x400520	2
0x4005e0	0x40052a	3
0x4005e8	0x40053c	4
0x4005f0	0x400533	5
0x4005f8	0x400533	6

# Code and Disassembled Executable

```
long switch_eg
(long x,
 long y,
 long z)
{
    long w = 1;
    switch(x) {
        case 1:
            w = y*z;
            break;
        case 2:
            w = y/z;
        case 3:
            w += z;
            break;
        case 5:
        case 6:
            w -= z;
            break;
        default:
            w = 2;
    }
    return w;
}
```

```
0000000000400508 <switch_eg>:
400508: 48 89 d1                mov     %rdx,%rcx
40050b: 48 83 ff 06             cmp     $0x6,%rdi
40050f: 77 2b                   ja      40053c
400511: ff 24 fd c8 05 40 00    jmpq    *0x4005c8(,%rdi,8)
400518: 48 89 f0                mov     %rsi,%rax
40051b: 48 0f af c2             imul    %rdx,%rax
40051f: c3                      retq
400520: 48 89 f0                mov     %rsi,%rax
400523: 48 99                   cqto
400525: 48 f7 f9                idiv    %rcx
400528: eb 05                   jmp     40052f
40052a: b8 01 00 00 00          mov     $0x1,%eax
40052f: 48 01 c8                add     %rcx,%rax
400532: c3                      retq
400533: b8 01 00 00 00          mov     $0x1,%eax
400538: 48 29 d0                sub     %rdx,%rax
40053b: c3                      retq
40053c: b8 02 00 00 00          mov     $0x2,%eax
400541: c3                      retq
400542: 66 2e 0f 1f 84 00 00    nopw    %cs:0x0(...)
400549: 00 00 00
40054c: 0f 1f 40 00             nopl    0x0(%rax)
```



# Disassembled Executable: Match w/ Code

```
long switch_eg
(long x, in rdi
 long y, in rsi
 long z) in rdx
{
    long w = 1;
    switch(x) {
        case 1:
            w = y*z;
            break;
        case 2:
            w = y/z;
        case 3:
            w += z;
            break;
        case 5:
        case 6:
            w -= z;
            break;
        default:
            w = 2;
    }
    return w;
}
```

```
0000000000400508 <switch_eg>:
400508: 48 89 d1                mov     %rdx,%rcx
40050b: 48 83 ff 06             cmp     $0x6,%rdi
40050f: 77 2b                   ja      40053c
400511: ff 24 fd c8 05 40 00    jmpq    *0x4005c8(,%rdi,8)
400518: 48 89 f0                mov     %rsi,%rax
40051b: 48 0f af c2             imul    %rdx,%rax
40051f: c3                      retq
400520: 48 89 f0                mov     %rsi,%rax
400523: 48 99                   cqto
400525: 48 f7 f9                idiv    %rcx
400528: eb 05                   jmp     40052f
40052a: b8 01 00 00 00          mov     $0x1,%eax
40052f: 48 01 c8                add     %rcx,%rax
400532: c3                      retq
400533: b8 01 00 00 00          mov     $0x1,%eax
400538: 48 29 d0                sub     %rdx,%rax
40053b: c3                      retq
40053c: b8 02 00 00 00          mov     $0x2,%eax
400541: c3                      retq
400542: 66 2e 0f 1f 84 00 00    nopw    %cs:0x0(...)
400549: 00 00 00
40054c: 0f 1f 40 00             nopl    0x0(%rax)
```

# Disassembled Executable: Match w/ Table

Jump table

Value	x
40053c	0
400518	1
400520	2
40052a	3
40053c	4
400533	5
400533	6

0000000000400508 <switch\_eg>:

```

400508: 48 89 d1      mov    %rdx,%rcx
40050b: 48 83 ff 06   cmp    $0x6,%rdi
40050f: 77 2b        ja     40053c
400511: ff 24 fd c8 05 40 00 jmpq   *0x4005c8(,%rdi,8)
400518: 48 89 f0      mov    %rsi,%rax
40051b: 48 0f af c2   imul   %rdx,%rax
40051f: c3          retq
400520: 48 89 f0      mov    %rsi,%rax
400523: 48 99        cqto
400525: 48 f7 f9     idiv   %rcx
400528: eb 05        jmp    40052f
40052a: b8 01 00 00 00 mov    $0x1,%eax
40052f: 48 01 c8     add    %rcx,%rax
400532: c3          retq
400533: b8 01 00 00 00 mov    $0x1,%eax
400538: 48 29 d0     sub    %rdx,%rax
40053b: c3          retq
40053c: b8 02 00 00 00 mov    $0x2,%eax
400541: c3          retq
400542: 66 2e 0f 1f 84 00 00 nopw   %cs:0x0(...)
400549: 00 00 00
40054c: 0f 1f 40 00  nopl   0x0(%rax)
    
```

# IA32 Switch Implementation

- Same general idea using 32-bit code
- Table entries 32 bits long (pointers)

```
long switch_eg(long x,    long y,    long z)
{
    long w = 1;
    switch(x) {
        ...
    }
    return w;
}
```

## Setup:

```
switch_eg:
    pushl    %ebp                # setup
    movl     %esp, %ebp         # setup
    movl     8(%ebp), %eax       # eax = x
    cmpl     $6, %eax           # compare x and 6
    ja       .L2                # if above (unsigned >) goto default
    jmp      *.L7(, %eax, 4)      # goto *jtab[x]
```

## Jump table

```
.section      .rodata
    .align 4
.L7:
    .long     .L2 # x = 0
    .long     .L3 # x = 1
    .long     .L4 # x = 2
    .long     .L5 # x = 3
    .long     .L2 # x = 4
    .long     .L6 # x = 5
    .long     .L6 # x = 6
```

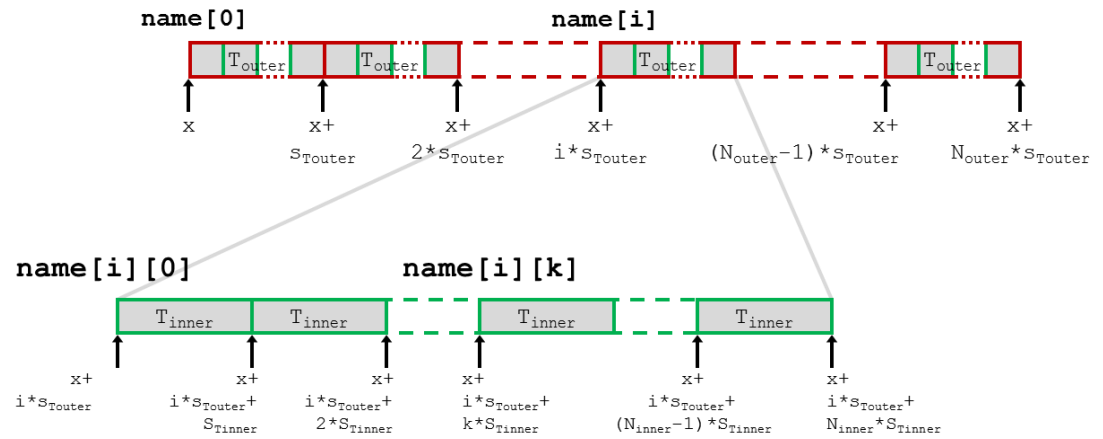
# Disassembled Executable (IA32)

```
8048422:  b8 02 00 00 00      mov     $0x2,%eax
8048427:  eb 2a               jmp     8048453 <switch_eg+0x43>
8048429:  b8 01 00 00 00      mov     $0x1,%eax
804842e:  66 90              xchg    %ax,%ax
8048430:  eb 14               jmp     8048446 <switch_eg+0x36>
8048432:  8b 45 10            mov     0x10(%ebp),%eax
8048435:  0f af 45 0c         imul    0xc(%ebp),%eax
8048439:  eb 18               jmp     8048453 <switch_eg+0x43>
804843b:  8b 55 0c            mov     0xc(%ebp),%edx
804843e:  89 d0              mov     %edx,%eax
8048440:  c1 fa 1f           sar     $0x1f,%edx
8048443:  f7 7d 10            idivl   0x10(%ebp)
8048446:  03 45 10            add     0x10(%ebp),%eax
8048449:  eb 08               jmp     8048453 <switch_eg+0x43>
804844b:  b8 01 00 00 00      mov     $0x1,%eax
8048450:  2b 45 10            sub     0x10(%ebp),%eax
8048453:  5d                  pop     %ebp
8048454:  c3                  ret
```

- single exit at bottom, code snippets jump to exit
- purpose of instruction at 0x804842e?

# Control Transfer Structures: Summary

- C Control
  - if-then-else
  - do-while
  - while, for
  - switch
- Assembler Control
  - Conditional jump
  - Conditional move
  - Indirect jump
  - Compiler generates code sequence to implement more complex control
- Standard Techniques
  - Loops converted to do-while form
  - Large switch statements use jump tables
  - Sparse switch statements may use decision trees



# Composite Data Structures

# Composite Data Structures

- **Arrays**
- Structures
- Unions

# Recap: Arrays

## ■ Declaration

$\langle T \rangle \text{ name} [\langle N \rangle]$

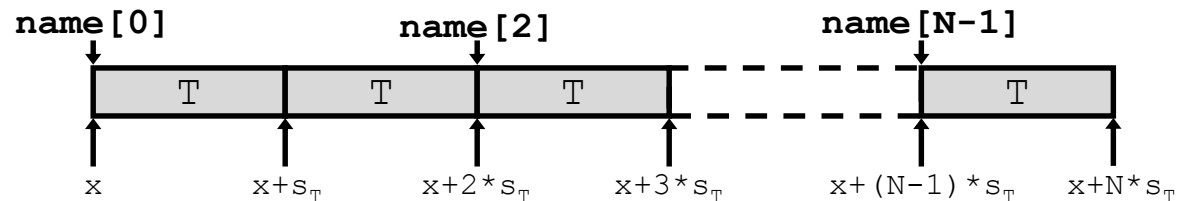
## ■ Size

- one element:
- entire array:

$$s_T = \text{sizeof}(T)$$

$$s_A = N * s_T$$

## ■ Memory layout



## ■ Address of i-th element

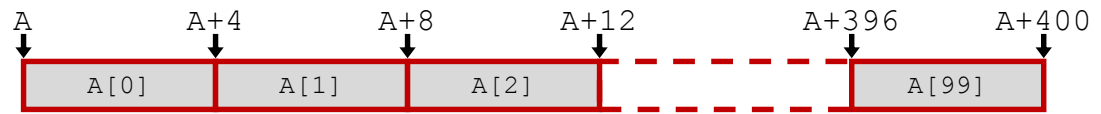
$$\text{adr}_i = \text{name} + i * s_T$$

## ■ Alignment

- array alignment = alignment of base type  $T$



# Example: 1-d Array



## One-dimensional array

```
#define N 100
```

```
int A[N];
```

```
int get(int i)
```

```
{  
    return A[i];  
}
```

```
int sum(void)
```

```
{  
    int i, sum = 0;
```

```
    for (i=0; i<N; i++) {  
        sum += A[i];  
    }
```

```
    return sum;
```

```
}
```

array1.c

```
$ gcc -m64 -O2 -S array1.c
```

```
get:
```

```
    movslq %edi, %rdi  
    movl   A(,%rdi,4), %eax  
    ret
```

```
sum:
```

```
    movl   $A, %edx  
    xorl   %eax, %eax  
.L4:  
    addl   (%rdx), %eax  
    addq   $4, %rdx  
    cmpq   $A+400, %rdx  
    jne    .L4  
    rep ret
```

array1.s

sum compiled as:

```
int sum(void)
```

```
{
```

```
    int *ap = A;
```

```
    int sum = 0;
```

```
    do {
```

```
        sum += *ap++;
```

```
    while (ap != A + 100);
```

```
    return sum;
```

```
}
```

# Recap: Multidimensional Arrays

- formed when `<type>` is an array type

$$\begin{aligned} & \langle T_{\text{outer}} \rangle \text{ name } [\langle N_{\text{outer}} \rangle] \\ \langle T_{\text{outer}} \rangle &= \langle T_{\text{inner}} \rangle \dots [\langle N_{\text{inner}} \rangle] \end{aligned}$$

combined notation

$$\langle T_{\text{inner}} \rangle \langle \text{name} \rangle [N_{\text{outer}}] [N_{\text{inner}}]$$

- **Size**

- one element:
- entire array:

$$s_{T_{\text{outer}}} = \text{sizeof}(T_{\text{outer}}) = N_{\text{inner}} * s_{T_{\text{inner}}}$$

$$s_{A_{\text{outer}}} = N_{\text{outer}} * s_{T_{\text{outer}}} = N_{\text{outer}} * N_{\text{inner}} * s_{T_{\text{inner}}}$$

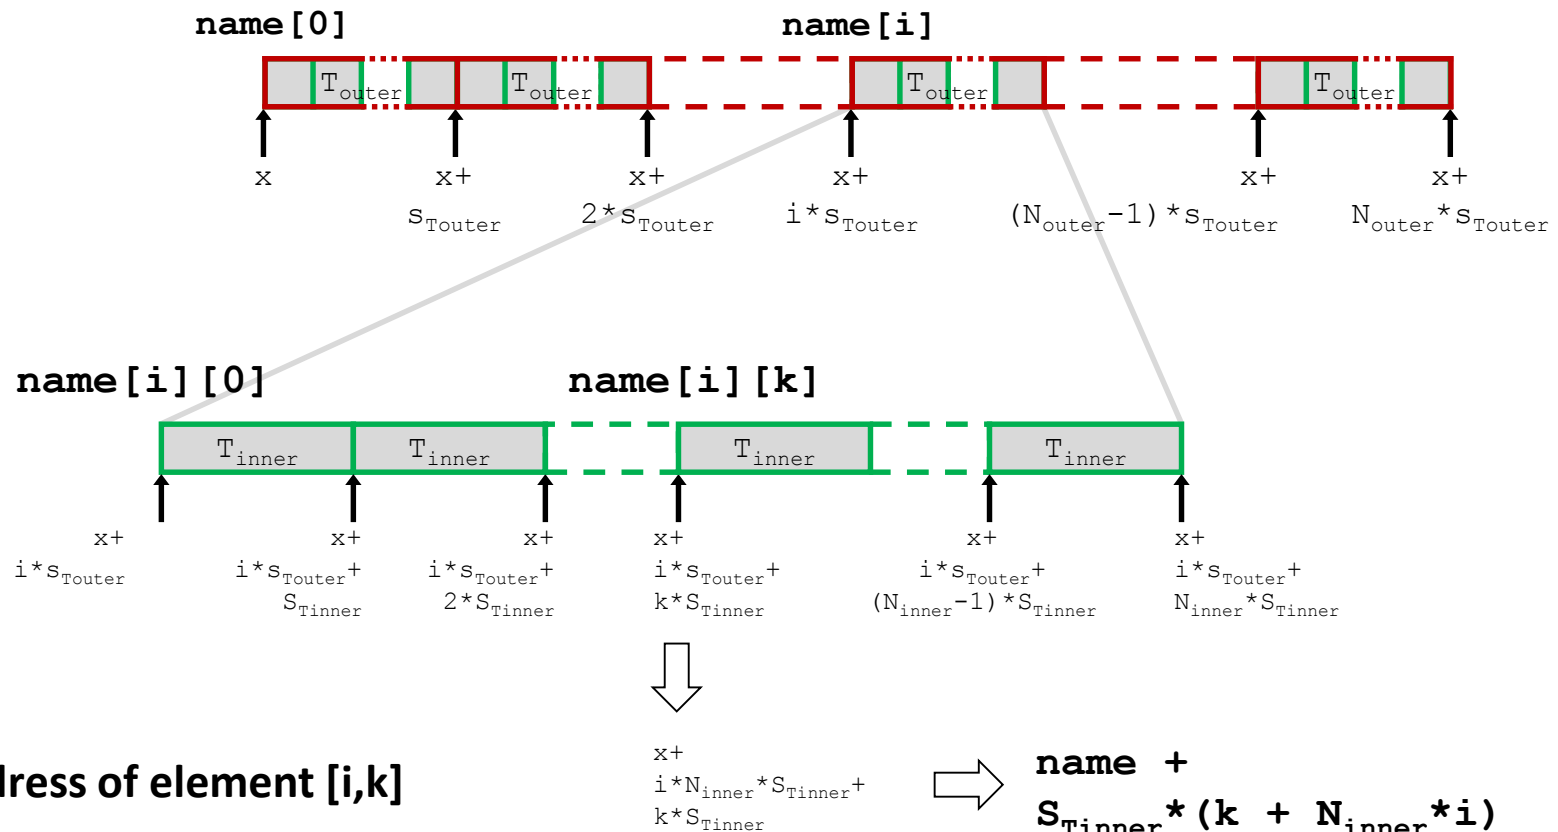
# Recap: Multidimensional Arrays

## ■ Memory layout (“row-major” layout)

$$\langle T_{\text{outer}} \rangle \text{ name } [\langle N_{\text{outer}} \rangle]$$

$$\langle T_{\text{outer}} \rangle = \langle T_{\text{inner}} \rangle \dots [\langle N_{\text{inner}} \rangle]$$

$$\langle T_{\text{inner}} \rangle \text{ } \langle \text{name} \rangle [N_{\text{outer}}] [N_{\text{inner}}]$$



## ■ Address of element [i,k]

- note how  $N_{\text{outer}}$  is not needed to compute `name[i][k]`

# Recap: Multidimensional Arrays

## ■ Generalization for n-dimensional array

$$\langle T_{\text{base}} \rangle \ \langle \text{name} \rangle [N_{D_n}] \dots [N_{D_2}] [N_{D_1}]$$

## ■ Size

- entire array

$$S_{D_n} = N_{D_n} * S_{T_{D_{n-1}}} = N_{D_n} * N_{D_{n-1}} * S_{T_{D_{n-2}}} = \dots = \mathbf{N_{D_n} * N_{D_{n-1}} * \dots * N_{D_1} * S_{T_{\text{base}}}}$$

- subdimension k ( $n \geq k \geq 1$ )

$$S_{D_k} = N_{D_k} * S_{T_{D_{k-1}}} = \dots = \mathbf{N_{D_k} * N_{D_{k-1}} * \dots * N_{D_1} * S_{T_{\text{base}}}}$$

# Recap: Multidimensional Arrays

## ■ Generalization for n-dimensional array

$$\langle T_{\text{base}} \rangle \ \langle \text{name} \rangle [N_{D_n}] \dots [N_{D_2}] [N_{D_1}]$$

## ■ Address

- $\text{name} [i_{D_n}] [i_{D_{n-1}}] \dots [i_{D_2}] [i_{D_1}]$

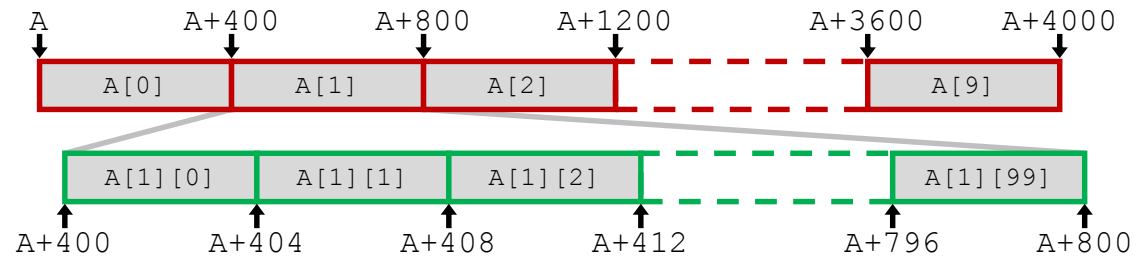
$$\text{name} + s_{T_{\text{base}}} * (i_{D_1} + N_{D_1} * (i_{D_2} + N_{D_2} * (\dots + N_{D_{n-1}} * i_{D_n}) \dots))$$

- again, note how  $N_{D_n}$  is not required for the address computation

→ this is why you can declare arrays with an open outermost dimension such as

```
int A[] [N] [K];
```

# Example: 2-d Array



## Two-dimensional array

```
#define N 10
#define M 100

int A[N][M];

int get(int i, int j)
{
    return A[i][j];
}

int sum(void)
{
    int i, j, sum = 0;

    for (i=0; i<N; i++) {
        for (j=0; j<M; j++) {
            sum += A[i][j];
        }
    }

    return sum;
}
```

array2.c

### get:

```
movslq %edi, %rdi
movslq %esi, %rsi
leaq    (%rdi,%rdi,4), %rax
leaq    (%rax,%rax,4), %rax
leaq    (%rsi,%rax,4), %rax
movl    A(%rax,4), %eax
ret
```

### sum:

```
movl    $A, %edx
movl    $A+4000, %esi
xorl    %eax, %eax
.L3:
    leaq    400(%rdx), %rcx
.L4:
    addl    (%rdx), %eax
    addq    $4, %rdx
    cmpq    %rcx, %rdx
    jne     .L4
    cmpq    %rdx, %rsi
    jne     .L3
    rep ret
```

array2.s

\$ gcc -m64 -O2 -S array2.c

# Example: 2-d Array

## Two-dimensional array (cont'd)

```
#define N 10
#define M 100

int A[N][M];

int sum(void)
{
    int i, j, sum = 0;

    for (i=0; i<N; i++) {
        for (j=0; j<M; j++) {
            sum += A[i][j];
        }
    }

    return sum;
}
```

array2.c

```
sum:
    movl    $A, %edx
    movl    $A+4000, %esi
    xorl    %eax, %eax
.L3:
    leaq    400(%rdx), %rcx
.L4:
    addl    (%rdx), %eax
    addq    $4, %rdx
    cmpq    %rcx, %rdx
    jne     .L4
    cmpq    %rdx, %rsi
    jne     .L3
    rep ret
```

array2.s

## sum compiled as:

```
int sum(void)
{
    int *Aptr = &A[0][0];
    int *endptr = &A[N][0];
    int sum = 0;

    do {
        int *endrow = Aptr + M;
        do {
            sum += *Aptr++;
        } while (Aptr < endrow);
    } while (Aptr < endptr);

    return sum;
}
```

# Example: Sum of a Column in a 2-D Array

Adding up a column in a 2-d array

```
#define N 100

int A[N][N];

int colsum(int c)
{
    int i, sum = 0;

    for (i=0; i<N; i++) {
        sum += A[i][c];
    }

    return sum;
}
```

array3.c

```
$ gcc -m64 -S -O2 array3.c
```

```
colsum:
    movslq %edi, %rcx
    xorl    %eax, %eax
    salq    $2, %rcx
    leaq    A(%rcx), %rdx
    addq    $A+40000, %rcx
.L2:
    addl    (%rdx), %eax
    addq    $400, %rdx
    cmpq    %rcx, %rdx
    jne     .L2
    rep ret
```

array3.s

colsum compiled as:

```
int colsum(int c)
{
    int sum = 0;
    int *ap = A+c;
    int *endp = A+10000+c;

    do {
        sum += *ap;
        ap = ap + 100;
    } while (ap != endp)

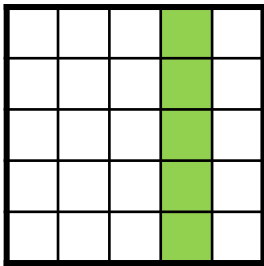
    return sum;
}
```



# Example: Sum of a Column in a 2-D Array

Adding up a column in a 2-d array

A (logical)



**colsum:**

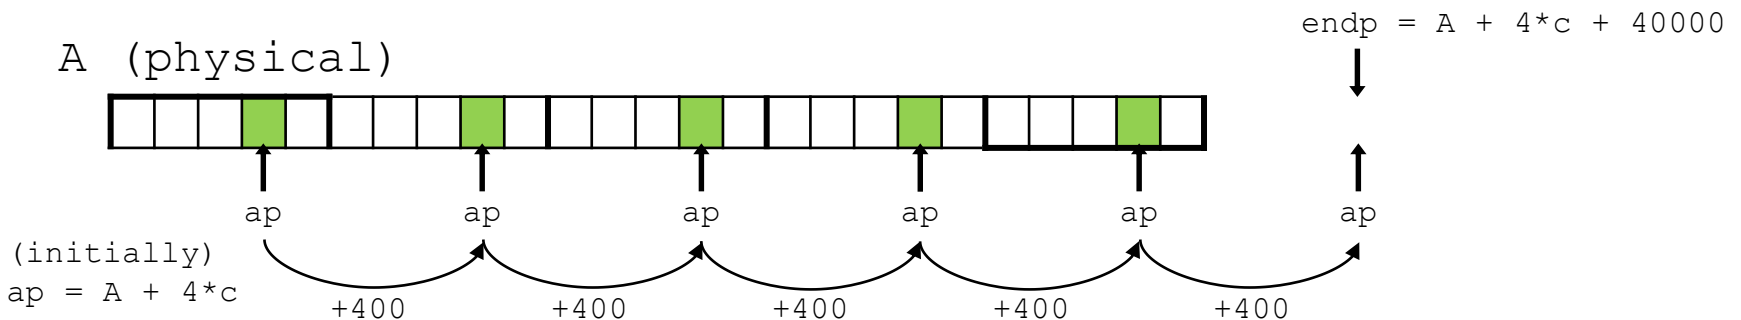
```
movslq %edi, %rcx
xorl    %eax, %eax
salq    $2, %rcx
leaq    A(%rcx), %rdx
addq    $A+40000, %rcx
```

**.L2:**

```
addl    (%rdx), %eax
addq    $400, %rdx
cmpq    %rcx, %rdx
jne     .L2
rep ret
```

array3.s

A (physical)



# Example: 4-d Array

## 4-dimensional array

```
#define N 64

int A[N][N][N][N];

int get(int i, int j, int k, int l)
{
    return A[i][j][k][l];
}

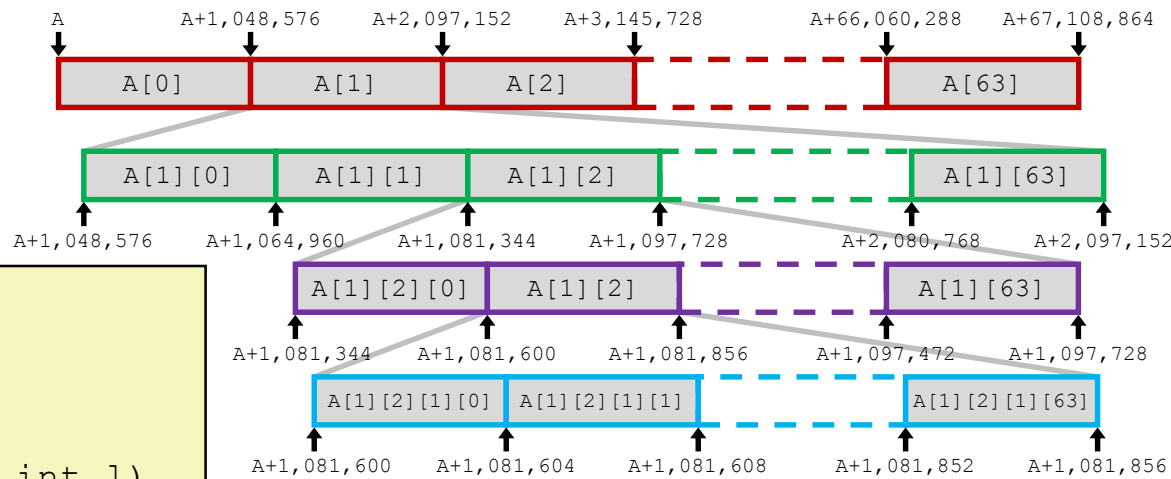
int sum(void)
{
    int i, j, k, l, sum = 0;

    for (i=0; i<N; i++)
        for (j=0; j<N; j++)
            for (k=0; k<N; k++)
                for (l=0; l<N; l++)
                    sum += A[i][j][k][l];

    return sum;
}
```

array4.c

\$ gcc -m64 -S -O2 array4.c



size of one element in:

1 <sup>st</sup> dimension	1,048,576	bytes
2 <sup>nd</sup> dimension	16,384	bytes
3 <sup>rd</sup> dimension	256	bytes
4 <sup>th</sup> dimension	4	bytes

get:

```
movslq %edi, %rdi
movslq %esi, %rsi
movslq %edx, %rdx
salq $6, %rdi
movslq %ecx, %rcx
addq %rdi, %rsi
salq $6, %rsi
addq %rsi, %rdx
salq $6, %rdx
addq %rcx, %rdx
movl A(,%rdx,4), %eax
ret
```

array4.s

# Example: 4-d Array

## 4-dimensional array

```
#define N 64
int A[N][N][N][N];
int sum(void)
{ int i, j, k, l, sum = 0;
  for (i=0; i<N; i++)
    for (j=0; j<N; j++)
      for (k=0; k<N; k++)
        for (l=0; l<N; l++)
          sum += A[i][j][k][l];
  return sum;
}
```

array4.c

```
int sum(void)
{   rdi = A+1,048,576+16,384;
    r8  = A+67,108,864+1,048,576+16384;
    sum (=eax) = 0;

13:  rsi = rdi - 1,048,576;
110: rdx = rsi - 16,384;
18:  rcx = rdx + 256;
16:  sum += *rdx++;
     if (rdx != rcx) goto 16;
     if (rdx != rsi) goto 18;
     rsi = rsi + 16,384;
     if (rsi != rdi) goto 110;
     rdi = rsi + 1,048,576;
     if (rdi != r8) goto 13;

    return sum;
}
```

gcc  
→

```
sum:
    movl    $A+1064960, %edi
    movl    $A+68173824, %r8d
    xorl    %eax, %eax
.L3:
    leaq    -1048576(%rdi), %rsi
.L10:
    leaq    -16384(%rsi), %rdx
.L8:
    leaq    256(%rdx), %rcx
.L6:
    addl    (%rdx), %eax
    addq    $4, %rdx
    cmpq    %rcx, %rdx
    jne     .L6
    cmpq    %rdx, %rsi
    jne     .L8
    addq    $16384, %rsi
    cmpq    %rdi, %rsi
    jne     .L10
    leaq    1048576(%rsi), %rdi
    cmpq    %r8, %rdi
    jne     .L3
    rep ret
```

array4.s

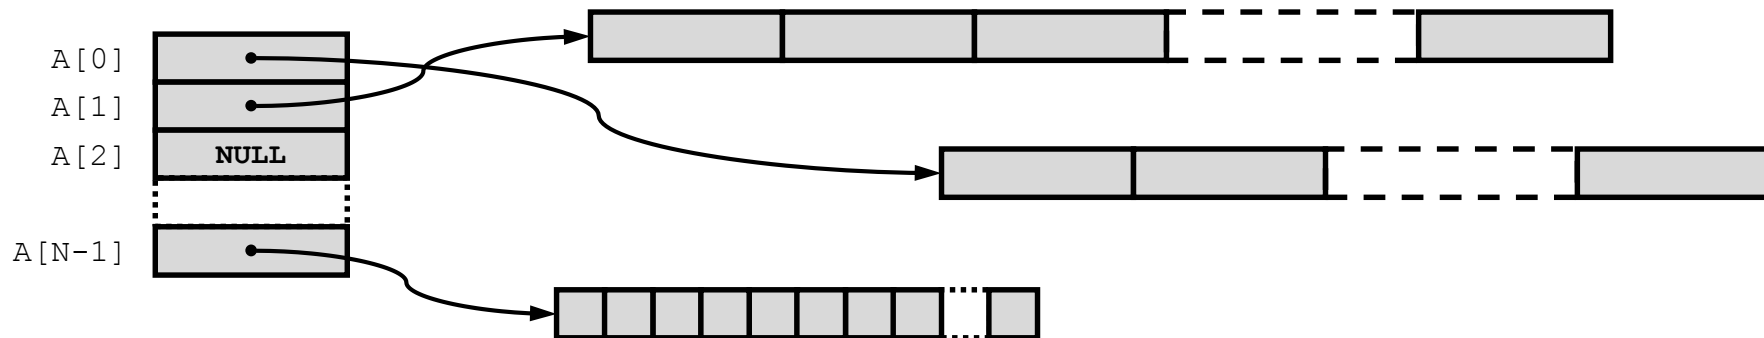
←  
asm-to-C

size of entire array	67,108,864	bytes
size of one element in:		
1 <sup>st</sup> dimension	1,048,576	bytes
2 <sup>nd</sup> dimension	16,384	bytes
3 <sup>rd</sup> dimension	256	bytes
4 <sup>th</sup> dimension	4	bytes

# Recap: Multilevel Arrays

- A multilevel array is simply an array of pointers to another array

`int *A[N]`



- pointed-to arrays can have different sizes
- elements in the pointer array can be NULL
- one extra memory access per pointer indirection
- both arrays, pointer array and pointed-to arrays can be multidimensional
- address calculation separate  
pointer array: element = pointer; pointed-to array: any type

# Example: Multidimensional vs. Multilevel Array

## Multidimensional (nested) array

```
int A[N][M];

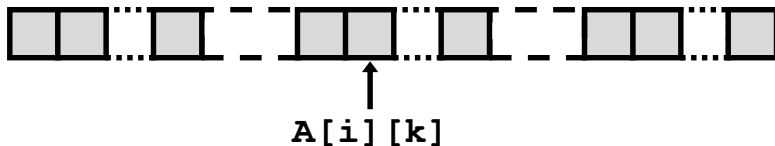
int get(int i, int k)
{
    return A[i][k];
}
```

## Multi-level array

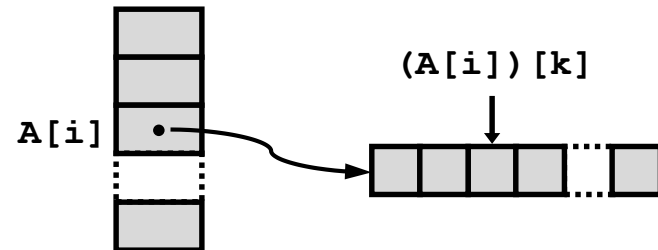
```
int *A[N];

int get(int i, int k)
{
    return A[i][k];
}
```

C code looks identical, but access very different:



**MEM[ A+4\* (k + M\*i) ]**



**MEM[ MEM[A+4\*i] + 4\*k ]**

# Example: Multidimensional vs. Multilevel Array

## Multidimensional (nested) array

```
#define N 5
#define M 10

long int A[N][M];

long int get(long int i, long int k)
{
    return A[i][k];
}
```

```
# parameters: %rdi = i, %rsi = k

<get>:
 0: lea    (%rdi,%rdi,4),%rax
 4: lea    (%rsi,%rax,2),%rax
 8: mov    A(,%rax,8),%rax
10: retq
```

**MEM[ A+8\*(k + 10\*i) ]**

## Multi-level array

```
#define N 5

long int *A[N];

long int get(long int i, long int k)
{
    return A[i][k];
}
```

```
# parameters: %rdi = i, %rsi = k

<get>:
 0: mov    A(,%rdi,8),%rax
 8: mov    (%rax,%rsi,8),%rax
 c: retq
```

**MEM[ MEM[A+8\*i] + 8\*k ]**

# Example: Multidimensional vs. Multilevel Array

## Multidimensional (nested) array

```
#define N 5

long int A[N][N][N];

long int get(long int i, long int k,
             long int l)
{
    return A[i][k][l];
}
```

```
# parameters: %rdi=i, %rsi=k, %rdx=l

<get>:
 0: lea    (%rdi,%rdi,4),%rax
 4: lea    (%rsi,%rsi,4),%rcx
 8: lea    (%rax,%rax,4),%rax
 c: add    %rcx,%rax
 f: add    %rax,%rdx
12: mov    A(,%rdx,8),%rax
1a: retq
```

## Multi-level array

```
#define N 5

long int **A[N];

long int get(long int i, long int k,
             long int l)
{
    return A[i][k][l];
}
```

```
# parameters: %rdi=i, %rsi=k, %rdx=l

<get>:
 0: mov    A(,%rdi,8),%rax
 8: mov    (%rax,%rsi,8),%rax
 c: mov    (%rax,%rdx,8),%rax
10: retq
```

**MEM[ A+8\*(1 + 5\*(k + 5\*i)) ]**

**MEM[ MEM[ MEM[A+8\*i] + 8\*k ] + 8\*l ]**

# Composite Data Structures

- Arrays
- **Structures**
- Unions



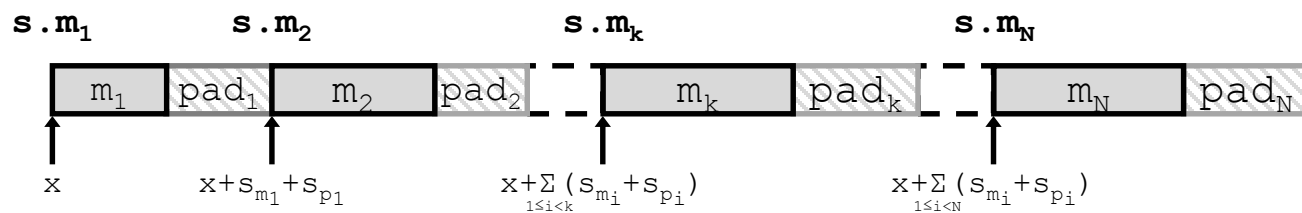
# Recap: Structures

## ■ Declaration

```
struct name {  
    <T1> <m1>;  
    <T2> <m2>;  
    ...  
    <TN> <mN>;  
};
```

## ■ Memory layout

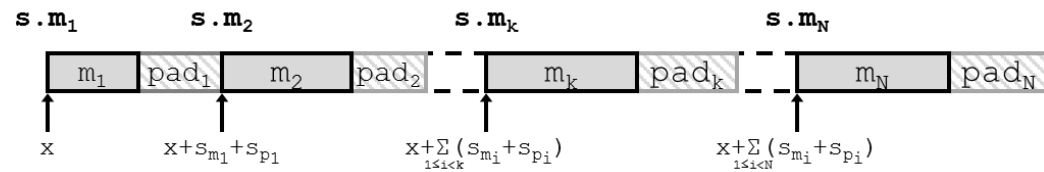
- consecutive memory region containing all members  $m_i$  in-order, non-overlapping, and properly aligned



## ■ Alignment

- struct alignment = maximum alignment requirement of any of its members
- member alignment = alignment requirement of member type
  - ▶ padding: “holes” in the memory layout to maintain alignment requirements (denoted  $pad_i$  above)

# Recap: Structures



## ■ Address of k-th member

- start of struct plus sum of sizes of all 1..k-1 members and paddings

$$\text{adr}_{m_k} = x + \sum_{1 \leq i < k} (s_{m_i} + s_{p_i})$$

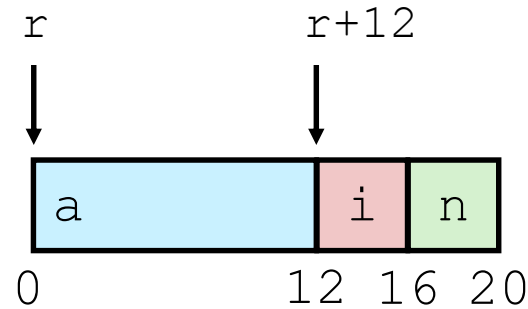
## ■ Size of struct

$$s_s = \sum_{1 \leq i \leq N} (s_{m_i} + s_{p_i})$$

- the last padding ( $pad_N$ ) is chosen such that the size of the struct is the next multiple of the biggest alignment requirement of any of its members
  - ▶ this comes in handy when declaring arrays of structs (all elements of the array will be automatically aligned)

# Example: Structure Access

```
struct rec {  
    int a[3];  
    int i;  
    struct rec *n;  
};
```



## ■ Accessing Structure Member

- Pointer indicates first byte of structure
- Access elements with offsets

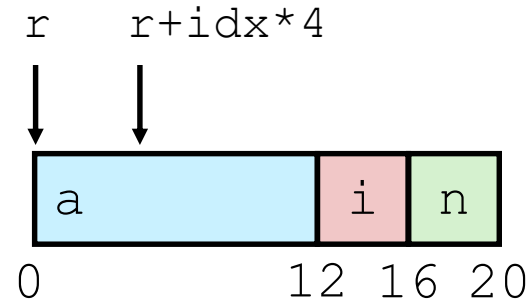
```
void  
set_i(struct rec *r,  
      int val)  
{  
    r->i = val;  
}
```

## x86\_64 assembly

```
# %edx = val  
# %rax = r  
movl %edx, 12(%rax)    # Mem[r+12] = val
```

# Example: Generating Pointer to Member

```
struct rec {  
    int a[3];  
    int i;  
    struct rec *n;  
};
```



## ■ Generating Pointer to Array Element

- Offset of each structure member determined at compile time
- Arguments
  - ▶ `%rdi`: `r`
  - ▶ `%rsi`: `idx`

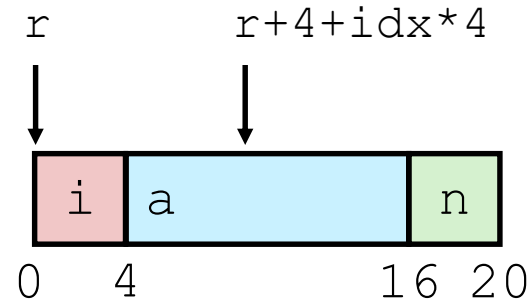
```
int *get_ap  
(struct rec *r, long int idx)  
{  
    return &r->a[idx];  
}
```

## x86-64 assembly

```
get_ap:  
    leaq    (%rdi,%rsi,4), %rax  
    ret
```

# Example: Generating Pointer to Member

```
struct rec {  
    int i;  
    int a[3];  
    struct rec *n;  
};
```



## ■ Generating Pointer to Array Element

- Offset of each structure member determined at compile time
- Arguments
  - ▶ `%rdi`: `r`
  - ▶ `%rsi`: `idx`

```
int *get_ap  
    (struct rec *r, long int idx)  
{  
    return &r->a[idx];  
}
```

## x86-64 assembly

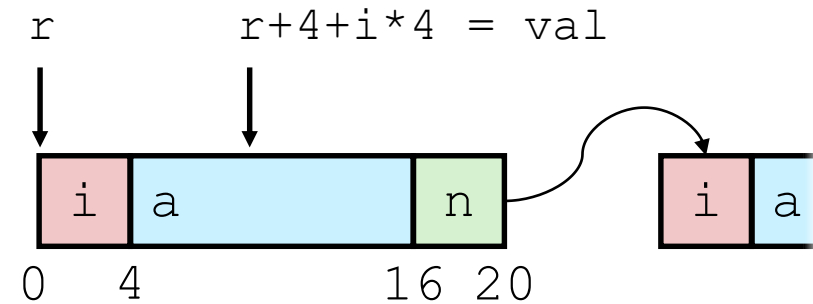
```
get_ap:  
    leaq    4(%rdi,%rsi,4), %rax  
    ret
```

# Example: Following a Linked List

```
#include <stdio.h>

struct rec {
    int i;
    int a[3];
    struct rec *n;
};

void set_val(struct rec *r, int val)
{
    while (r != NULL) {
        int i = r->i;
        r->a[i] = val;
        r = r->n;
    }
}
```



- Set element `r->a[r->i]` to `val` in entire linked list

- Arguments

- ▶ `%rdi`: `r`
- ▶ `%rsi(%esi)`: `val`

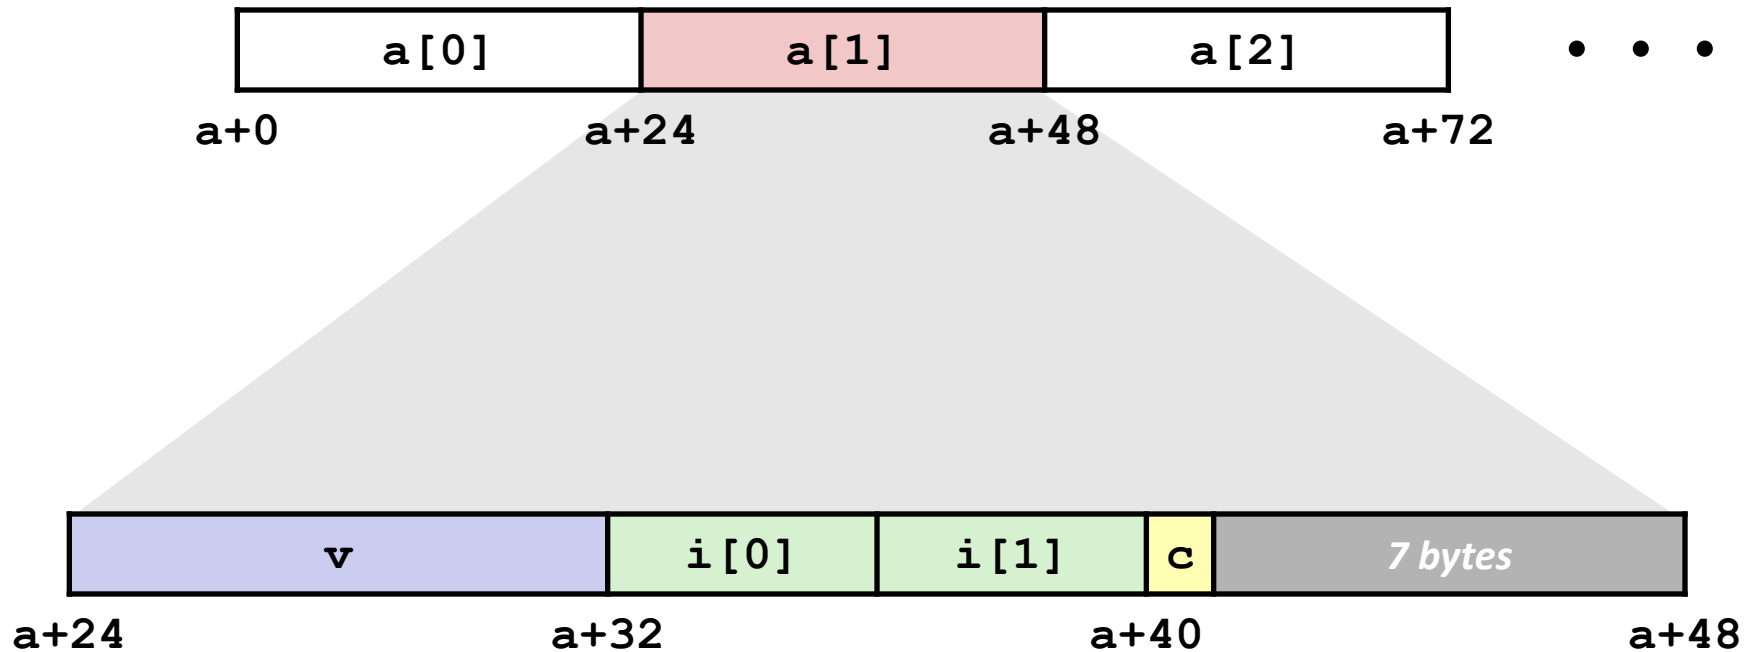
## x86-64 assembly

```
set_val:
    jmp     .L7
.L5:
    movslq  (%rdi), %rax           # r->i
    movl    %esi, 4(%rdi,%rax,4)   # r->a[r->i]
    movq    16(%rdi), %rdi         # r->next
.L7:
    testq   %rdi, %rdi            # r
    jne     .L5
    rep ret
```

# Example: Arrays of Structures

- Overall structure length multiple of K
- Satisfy alignment requirement for every element

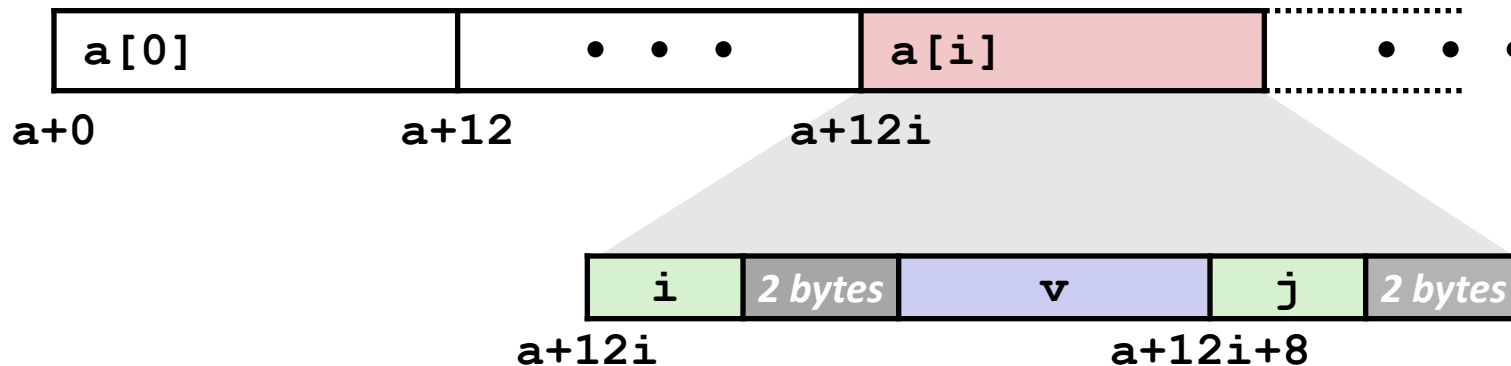
```
struct S2 {  
    double v;  
    int i[2];  
    char c;  
} a[10];
```



# Example: Accessing Array Elements

- Compute array offset  $12i$ 
  - `sizeof(S3)`, including alignment spacers
- Element  $j$  is at offset 8 within structure
- Assembler gives offset  $a+8$ 
  - Resolved during linking

```
struct S3 {  
    short i;  
    float v;  
    short j;  
} a[10];
```



```
short get_j(int idx)  
{  
    return a[idx].j;  
}
```

```
# %eax = idx  
leal (%eax,%eax,2),%eax # 3*idx  
movswl a+8(,%eax,4),%eax
```



# Composite Data Structures

- Data Alignment
- Arrays
- Structures
- **Unions**

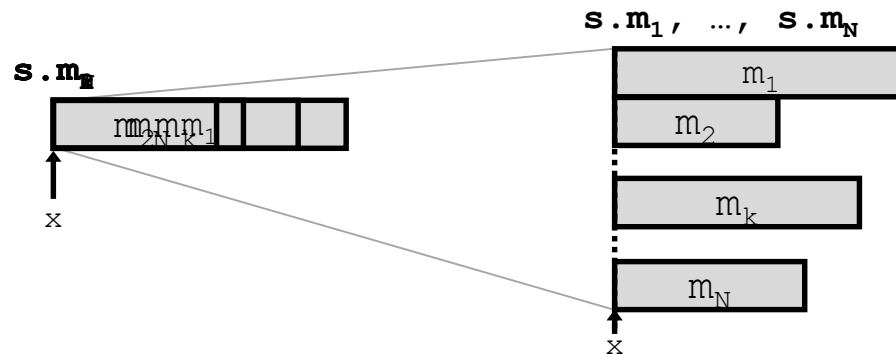
# Recap: Unions

## ■ Declaration

```
union name {  
    <T1> <m1>;  
    <T2> <m2>;  
    ...  
    <TN> <mN>;  
};
```

## ■ Memory layout

- consecutive memory region containing all members  $m_i$ , and properly aligned
- all members *are located at offset 0 and overlap in memory*



# Recap: Unions

## ■ Alignment

- union alignment = maximum alignment requirement of any of its members
- member alignment = alignment requirement of member type

## ■ Address of k-th member

- start of union

$$\text{adr}_{m_k} = x$$

## ■ Size of union

$$s_u = \max_{1 \leq i \leq N} (s_{m_i})$$

## Example: Using Unions to Access Bit Patterns

- Task: print bit pattern of a floating point number
- Try 1:

```
#include <stdio.h>

unsigned int get_bitpattern(float f)
{
    return (unsigned int)f;
}

void main(void)
{
    float f = 3.14159265358979323846;
    unsigned int u, i;

    u = get_bitpattern(f);

    printf("%12.10f = ", f);
    for (i=sizeof(u)*8; i>0; i--) {
        printf("%c", (u & (1 << (i-1)) ? '1' : '0'));
    }
    printf("b = 0x%08x\n", u);
}
```

Output looks suspiciously wrong:

[illegible]

## The assembly reveals

```
get_bitpattern:
    subl    $20, %esp
    fnstcw  14(%esp)
    movzwl  14(%esp), %eax
    flds    24(%esp)
    movb    $12, %ah
    movw    %ax, 12(%esp)
    fldcw   12(%esp)
    fistpq  (%esp)      # convert float to int
    fldcw   14(%esp)
    movl    (%esp), %eax
    addl    $20, %esp
    ret
```

indeed, this is a conversion float  $\rightarrow$  int.

# Example: Using Unions to Access Bit Patterns

## ■ Try 2:

```
#include <stdio.h>

unsigned int get_bitpattern(float f)
{
    union {
        float f;
        unsigned int u;
    } fu;
    fu.f = f;
    return fu.u;
}

void main(void)
{
    float f = 3.14159265358979323846;
    unsigned int u, i;

    u = get_bitpattern(f);

    printf("%12.10f = ", f);
    for (i=sizeof(u)*8; i>0; i--) {
        printf("%c", (u & (1 << (i-1)) ? '1' : '0'));
    }
    printf("b = 0x%08x\n", u);
}
```

convert2.c

This is what we want

```
get_bitpattern:
    movl    4(%esp), %eax
    ret
```

Output:

```
$ gcc -m32 -O3 -o convert2 convert2.c
$ ./convert2
3.1415927410 = 01000000010010010000111111011011b = 0x40490fdb
```

# Calling Convention

# Calling Convention on x86\_64

- **Arguments passed to functions via registers**
  - If more than 6 integral parameters, then pass rest on stack
  - These registers can be used as caller-saved as well
- **All references to stack frame via stack pointer `%rsp`**
- **Register saving convention**
  - 6 “callee saved”
  - 2 “caller saved”
  - 1 return value (also usable as caller saved)
  - 1 special (stack pointer)

# Calling Convention on x86\_64

%rax	Caller saved / Return value
%rbx	Callee saved
%rcx	Caller saved / Argument #4
%rdx	Caller saved / Argument #3
%rsi	Caller saved / Argument #2
%rdi	Caller saved / Argument #1
%rsp	Stack pointer
%rbp	Callee saved

%r8	Caller saved / Argument #5
%r9	Caller saved / Argument #6
%r10	Caller saved / Caller saved
%r11	Caller Saved
%r12	Callee saved
%r13	Callee saved
%r14	Callee saved
%r15	Callee saved



# Calling Convention on IA32

- Arguments passed to functions via stack
- References to stack frame via base pointer `%ebp`
- Register saving convention
  - 3 “callee saved”
  - 3 “caller saved”
  - 1 return value (also usable as caller saved)
  - 2 special (stack pointer)

# IA32/Linux+Windows Calling Convention

## ■ **%eax, %ecx, %edx**

- caller saved prior to call (if values are used later)
- %eax used to return integer value

## ■ **%ebx, %esi, %edi**

- callee saved (if used)

## ■ **%esp, %ebp**

- used to manage the stack frames
- must restore original values upon exit from procedure (= special form of callee saved)

%eax	Caller saved / Return value
%ecx	Caller saved
%edx	Caller saved
%ebx	Callee saved
%esi	Callee saved
%edi	Callee saved
%esp	Stack pointer
%ebp	Frame pointer

# Implementation of Linux/x86\_64 System Calls

- System calls are invoked via a special `syscall/sysenter` instruction
  - parameters passed through registers (not the stack)
    - ▶ `rax`: system call number
    - ▶ `rdi, rsi, rdx, rcx, r8, r9`: (up to) 6 arbitrary arguments; interpretation depends on the invoked system call.
      - the kernel uses a different set of registers:  
`rdi, rsi, rdx, r10, r8, r9`
    - ▶ result of `syscall` placed in `rax`
      - `>0`: system call succeeded, result interpretation depends on `syscall`
      - `<0`: error
- `libc` provides convenient wrappers for most system calls

# Implementation of Linux/x86\_64 System Calls

## ■ Directly invoking Linux/x86\_64 System calls

```
#include <stdlib.h>
#include <unistd.h>

int main(void)
{
    write(1, "hello, world\n", 13);
    exit(0);
}                                     hello.c
```

```
$ gcc -o hello64 hello.c
$ ./hello64
hello, world
$
```

```
int main(void)
{
    my_write(1, "hello, world\n", 13);
    my_exit(0);
}                                     myhello.c
```

```
.globl    my_write
.globl    my_exit

.section .text

# ssize_t my_write(int fd, const void *string,
#                                     size_t count)
#
# parameters on function entry:
#   rdi: fd
#   rsi: buf
#   rdx: count

my_write:

    mov    $1, %rax                # rax: 1 (write syscall ID)
                                    # rdi: fd
                                    # rsi: buf
                                    # rdx: count

    syscall                        # invoke system call
                                    # result in rax

    ret

...                                mylib64.s
```

# Implementation of Linux/x86\_64 System Calls

## ■ Directly invoking Linux/x86\_64 System calls

```
int main(void)
{
    my_write(1, "hello, world\n", 13);
    my_exit(0);
}                                     myhello.c
```

```
...
# void my_exit(int status)
#
# parameters on function entry:
#   rdi: status
my_exit:

    mov    $60, %eax                # rax: 60 (exit syscall ID)
                                      # rdi: status

    syscall                        # invoke system call

    # exit does not return, but just to be sure

    ret                               mylib64.s
```

```
$ gcc -o myhello64 myhello.c mylib64.s
$ ./myhello64
hello, world
$
```

# Implementation of Linux/IA32 System Calls

- Application programs use system calls to request a service from the kernel
- System calls are invoked via the `int 0x80` instruction
  - parameters passed through registers (not the stack)
    - ▶ `eax`: system call number
    - ▶ `ebx, ecx, edx, esi, edi, ebp`: (up to) 6 arbitrary arguments (interpretation depends on the invoked system call)
    - ▶ `esp` cannot be used (overwritten when the CPU enters kernel mode)
- `libc` provides convenient wrappers for most system calls

# Implementation of Linux/IA32 System Calls

## ■ Directly invoking Linux/IA32 System calls

```
#include <stdlib.h>
#include <unistd.h>

int main(void)
{
    write(1, "hello, world\n", 13);
    exit(0);
}                                     hello.c
```

```
$ gcc -m32 -o hello32 hello.c
$ ./hello32
hello, world
$
```

```
int main(void)
{
    my_write(1, "hello, world\n", 13);
    my_exit(0);
}                                     myhello.c
```

```
.globl my_write
.globl my_exit

.section .text

# int my_write(int fd, const void *buf,
               unsigned int count);
my_write:

    push    %ebp
    mov     %esp, %ebp
    push    %ebx

    mov     $4, %eax           # eax: 4 (write syscall)
    mov     8(%ebp), %ebx      # ebx: fd      (ebp+8)
    mov     12(%ebp), %ecx     # ecx: buf      (ebp+12)
    mov     16(%ebp), %edx     # edx: count   (ebp+16)

    int     $0x80             # invoke system call
                                # result in eax

    pop     %ebx
    pop     %ebp

    ret

...                               mylib32.s
```

# Implementation of Linux/IA32 System Calls

## ■ Directly invoking Linux/IA32 System calls

```
int main(void)
{
    my_write(1, "hello, world\n", 13);
    my_exit(0);
}                                     myhello.c
```

```
...
# void my_exit(int nr)
my_exit:

    push    %ebp
    mov     %esp, %ebp
    push    %ebx

    mov     $1, %eax                # eax: 1 (exit syscall)
    mov     8(%ebp), %ebx           # ebx: nr (ebp+8)

    int     $0x80                  # invoke system call

    pop     %ebx
    pop     %ebp

    ret                                     mylib32.s
```

```
$ gcc -m32 -o myhello32 myhello.c mylib32.s
$ ./myhello32
hello, world
$
```



# Linux Kernel System Call Implementation

## ■ Linux kernel system call table

- in the Linux kernel source directory, see  
arch/x86/entry/syscalls/

syscall_64.tbl	for 64-bit system calls
syscall_32.tbl	for 32-bit system calls

## ■ User-level SYSCALL invocation

- glibc provides wrappers for each system call (“write” → SYSCALL in assembly)
- source builds syscalls dynamically; easier to look at compiled libc.a

```
/tmp $ ar t /usr/lib/libc.a
/tmp $ ar x /usr/lib/libc.a write.o
/tmp $ objdump -d write.o
0000000000000000 <__libc_write>:
...
    6d:    b8 01 00 00 00      mov     $0x1,%eax
    72:    0f 0f               syscall
...
```

# Linux Kernel System Call Implementation

## ■ Kernel SYSCALL entry point

- `ENTRY(entry_SYSCALL_64)` in `arch/x86/entry/entry64.S`
- depending on kernel version less or more complicated

## ■ Kernel implementation of a system call

- `grep` kernel tree for `"SYSCALL_DEFINE.\?(<syscall>,"`
- for example, to find the `write` system call:

```
/usr/src/linux $ grep -rA3 "SYSCALL_DEFINE.\?(write," *
SYSCALL_DEFINE3(write, unsigned int, fd, const char __user *,
                buf, size_t, count)
{
    struct fd f = fdget_pos(fd);
```

```
absdiff:
    cmpq    %rsi, %rdi
    jg      .L4
    movq    %rsi, %rax
    subq    %rdi, %rax
    jmp     .L7
.L4:
    movq    %rdi, %rax
    subq    %rsi, %rax
.L7:
    ret
```

# Module Summary

# Translation of C to Assembly

## ■ Control flow constructs

- No high-level control flow constructs at the assembly level
- All control flow is altered using a combination of comparison and jump instructions
- for → while → do while → assembly
- (Larger) switch statements use efficient jump tables

## ■ Composite data structures

- No composite data structures at the assembly level
- Translated by the compiler into offsets into a linear memory space

## ■ Calling conventions

- Caller vs callee-saved registers
- Parameter passing
  - ▶ via the stack on IA32, in registers on x86\_64