

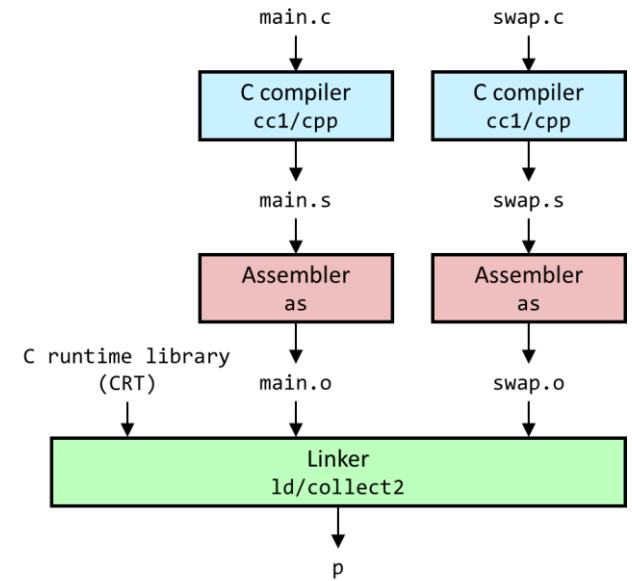
REVE1

Program Execution

```
1: 48 89 e5                mov    %rsp,%rbp
4: 48 c7 05 00 00 00 00    movq   $0x0,0x0(%rip)
b: 00 00 00 00
   7: R_X86_64_PC32 .bss-0x8
   b: R_X86_64_32S buf+0x4
f: 48 8b 05 00 00 00 00    mov     0x0(%rip),%rax
   12: R_X86_64_PC32 bufp0-0x4
16: 8b 00                mov     (%rax),%eax
18: 89 45 fc                mov     %eax,-0x4(%rbp)
1b: 48 8b 05 00 00 00 00    mov     0x0(%rip),%rax
   1e: R_X86_64_PC32 bufp0-0x4
22: 48 8b 15 00 00 00 00    mov     0x0(%rip),%rdx
   25: R_X86_64_PC32 .bss-0x4
29: 8b 12                mov     (%rdx),%edx
```

Module Outline

- The Life Cycle of a Program
- Linking Overview
- Executable and Linkable Format
- Linking
 - Symbol Resolution
 - Symbol Relocation
- Module Summary



Life Cycle of a Program

A Simple C Program

```
int buf[2] = {1, 2};

int main() {
    printf("[ %d, %d ]\n",
           buf[0],buf[1]);
    swap();
    printf("[ %d, %d ]\n",
           buf[0],buf[1]);
    return EXIT_SUCCESS;
}                                     main.c
```

```
extern int buf[];

int *bufp0 = &buf[0];
static int *bufp1;

void swap() {
    int temp;

    bufp1 = &buf[1];
    temp = *bufp0;
    *bufp0 = *bufp1;
    *bufp1 = temp;
}                                     swap.c
```

- Compile to executable program with GCC

```
$ gcc -O2 -Wall -g -o p main.c swap.c
```

```
main.c: In function 'main':
```

```
main.c:9:3: warning: implicit declaration of function 'swap'
```

```
    9 |     swap();
      |     ^~~~
```

```
$ ./p
```

```
[ 1, 2 ]
```

```
[ 2, 1 ]
```

```
$
```

A Simple C Program

```
int buf[2] = {1, 2};

int main() {
    printf(...);
    swap();
    printf(...);
}
```

main.c

```
extern int buf[];

int *bufp0 = &buf[0];
static int *bufp1;

void swap() {
    int temp;

    bufp1 = &buf[1];
    temp = *bufp0;
    *bufp0 = *bufp1;
    *bufp1 = temp;
}
```

swap.c

- **GCC is, in fact, a *compiler driver***
- Executes a series of commands that transform the source code into an executable

```
$ gcc -v -O2 -Wall -g -o p main.c swap.c
Using built-in specs.
COLLECT_GCC=gcc
COLLECT_LTO_WRAPPER=/usr/libexec/gcc/x86_64-pc-linux-gnu/11.3.0/lto-wrapper
Target: x86_64-pc-linux-gnu
Configured with: /var/tmp/portage/sys-devel/gcc-11.3.0/work/gcc-11.3.0/configure
--host=x86_64-pc-linux-gnu --build=x86_64-pc-linux-gnu --prefix=/usr
--bindir=/usr/x86_64-pc-linux-gnu/gcc-bin/11.3.0 --includedir=...
Thread model: posix
...
```

A Simple C Program

■ Making sense of GCC's compilation steps

```
int buf[2] = {1, 2};

int main() {
    printf(...);
    swap();
    printf(...);
    return 0;
}                                     main.c
```

```
extern int buf[];

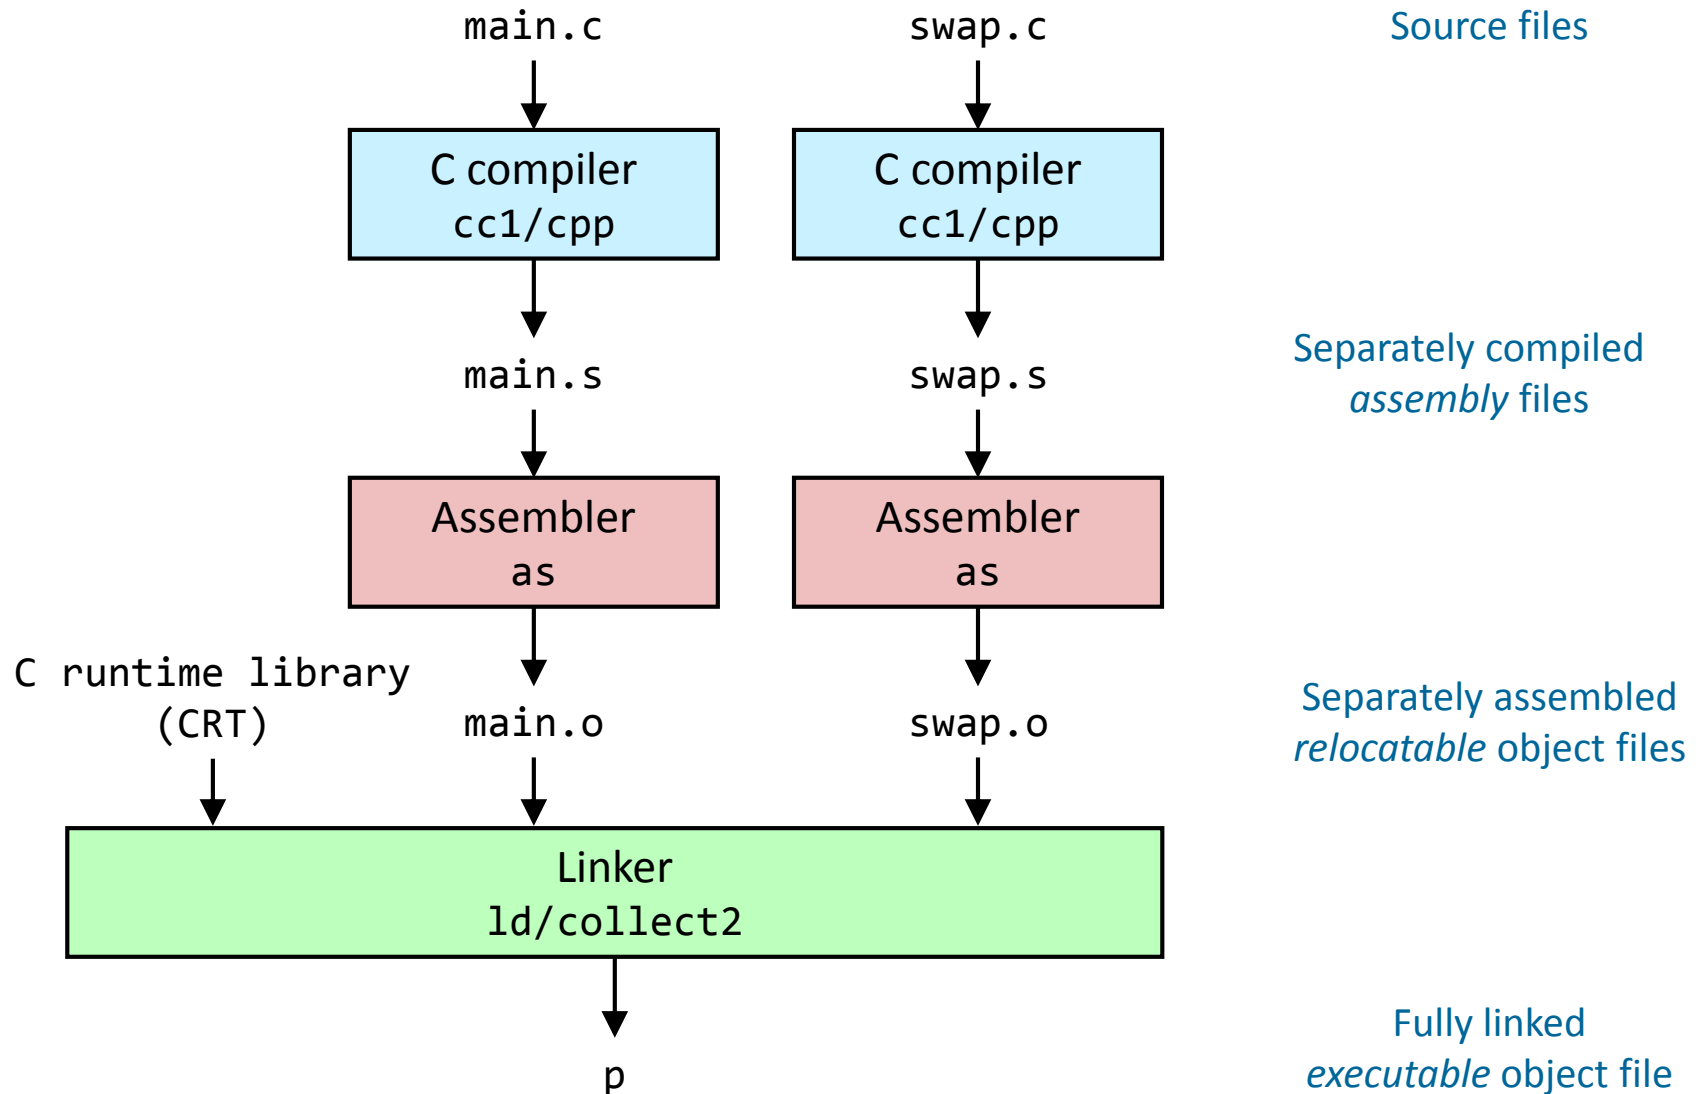
int *bufp0 = &buf[0];
static int *bufp1;

void swap() {
    int temp;

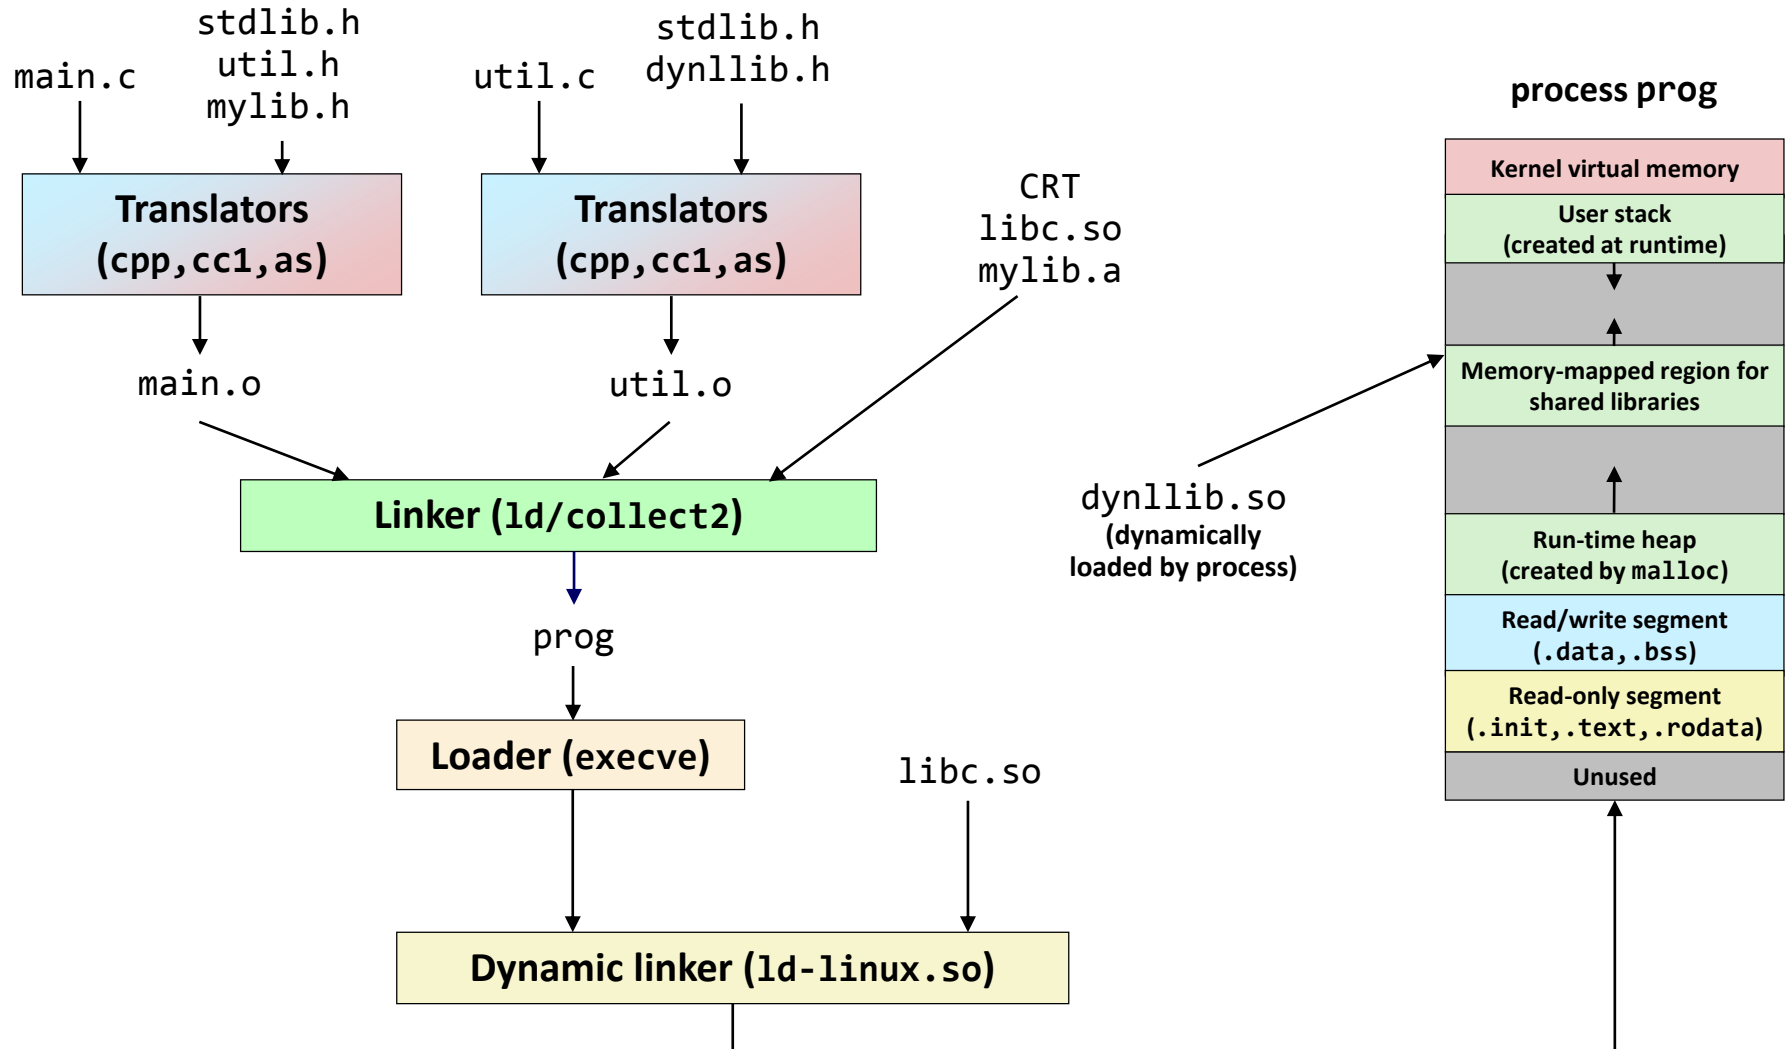
    bufp1 = &buf[1];
    temp = *bufp0;
    *bufp0 = *bufp1;
    *bufp1 = temp;
}                                     swap.c
```

```
$ ( gcc -v -O2 -Wall -g -o p main.c swap.c 2>&1 ) | \
  grep "^ /usr/lib" | grep -v include | \
  grep -E "^ /usr[^\ ]*"
/usr/libexec/gcc/x86_64-pc-linux-gnu/11.3.0/cc1 -quiet -v main.c -quiet -dumpdir p- -dumpbase
main.c -dumpbase-ext .c -mtune=generic -march=x86-64 -g -O2 -Wall -version -o /tmp/ccUMd0qR.s
/usr/lib/gcc/x86_64-pc-linux-gnu/11.3.0/../../../../x86_64-pc-linux-gnu/bin/as -v --gdwarf-5 --64
-o /tmp/ccFGruA6.o /tmp/ccUMd0qR.s
/usr/libexec/gcc/x86_64-pc-linux-gnu/11.3.0/cc1 -quiet -v swap.c -quiet -dumpdir p- -dumpbase
swap.c -dumpbase-ext .c -mtune=generic -march=x86-64 -g -O2 -Wall -version -o /tmp/ccUMd0qR.s
/usr/lib/gcc/x86_64-pc-linux-gnu/11.3.0/../../../../x86_64-pc-linux-gnu/bin/as -v --gdwarf-5 --64
-o /tmp/ccyXZYrW.o /tmp/ccUMd0qR.s
/usr/libexec/gcc/x86_64-pc-linux-gnu/11.3.0/collect2 -plugin /usr/libexec/gcc/x86_64-pc-linux-
gnu/11.3.0/liblto_plugin.so -plugin-opt=/usr/libexec/gcc/x86_64-pc-linux-gnu/11.3.0/lto-wrapper -
plugin-opt=-fresolution=/tmp/ccU0tDGc.res -plugin-opt=-pass-through=-lgcc -plugin-opt=-pass-
through=-lgcc_s -plugin-opt=-pass-through=-lc -plugin-opt=-pass-through=-lgcc -plugin-opt=-pass-
through=-lgcc_s --eh-frame-hdr -m elf_x86_64 -dynamic-linker /lib64/ld-linux-x86-64.so.2 -pie -o p
/usr/lib/gcc/x86_64-pc-linux-gnu/11.3.0/../../../../lib64/Scrt1.o /usr/lib/gcc/x86_64-pc-linux-
gnu/11.3.0/../../../../lib64/crti.o /usr/lib/gcc/x86_64-pc-linux-gnu/11.3.0/crtbeginS.o -
L/usr/lib/gcc/x86_64-pc-linux-gnu/11.3.0 -L/usr/lib/gcc/x86_64-pc-linux-
gnu/11.3.0/../../../../lib64 -L/lib/./lib64 -L/usr/lib/./lib64 -L/usr/lib/gcc/x86_64-pc-linux-
gnu/11.3.0/../../../../x86_64-pc-linux-gnu/lib -L/usr/lib/gcc/x86_64-pc-linux-gnu/11.3.0/../../../../
/tmp/ccFGruA6.o /tmp/ccyXZYrW.o -lgcc --push-state --as-needed -lgcc_s --pop-state -lc -lgcc --
push-state --as-needed -lgcc_s --pop-state /usr/lib/gcc/x86_64-pc-linux-gnu/11.3.0/crtendS.o
/usr/lib/gcc/x86_64-pc-linux-gnu/11.3.0/../../../../lib64/crtn.o
```

A Simple C Program



From Source Code to a Running Process



Problems to Solve

■ Separate compilation of individual C files

- How does the compiler know about functions and variables defined in other files to make sure the type matches?
- How can we generate addresses to call functions / access variables if we do not know where they are located in memory?
- How does the system know which dynamic libraries to load when executing a binary?
- How does the application know how to call functions from dynamic libraries?

```
$ gcc -c main.c
$ readelf -s main.o
```

```
...
```

Num:	Size	Type	Bind	Ndx	Name
9:	8	OBJECT	GLOBAL	3	buf
10:	83	FUNC	GLOBAL	1	main
13:	0	NOTYPE	GLOBAL	UND	swap

Linking Overview

Why Linkers?

■ Reason 1: Modularity

- Program can be written as a collection of smaller source files, rather than one monolithic mass.
- Can build libraries of common functions (more on this later)
 - ▶ e.g., Math library, standard C library

Why Linkers? (cont)

■ Reason 2: Efficiency

- Time: Separate compilation
 - ▶ Change one source file, compile, and then relink.
 - ▶ No need to recompile other source files.
- Space: Libraries
 - ▶ Common functions can be aggregated into a single file...
 - ▶ yet executable files and running memory images contain only code for the functions they actually use.

What Do Linkers Do?

■ Step 1. Symbol resolution

- Programs define and reference symbols (variables and functions):
 - ▶ `void swap() {...}` `/* define symbol swap */`
 - ▶ `swap();` `/* reference symbol swap */`
 - ▶ `int *bufp0 = &buf[0];` `/* define symbol bufp0, reference buf */`
- Symbol definitions are stored (by the compiler) in a symbol table.
 - ▶ Symbol table is an array of structs
 - ▶ Each entry includes name, size, and location of symbol
- Linker associates each symbol reference with exactly one symbol definition.

What Do Linkers Do? (cont)

■ Step 2. Relocation

- Merges separate code and data sections into single sections
- Relocates symbols from their relative locations in the .o files to their final absolute memory locations in the executable.
- Updates all references to these symbols to reflect their new positions.

ELF header
Segment header table (required for executables)
.text section
.rodata section
.data section
.bss section
.symtab section
.rel.txt section
.rel.data section
.debug/.line section
Section header table

Executable and Linkable Format (ELF)

Three Kinds of Object Files (Modules)

■ Relocatable object file (.o file)

- Contains code and data in a form that can be combined with other relocatable object files to form executable object file.
 - ▶ Each .o file is produced from exactly one source (.c) file

■ Executable object file (a.out file)

- Contains code and data in a form that can be copied directly into memory and then executed.

■ Shared object file (.so file)

- Special type of relocatable object file that can be loaded into memory and linked dynamically, at either load time or run-time.
- Called Dynamic Link Libraries (DLLs) by Windows

Executable and Linkable Format (ELF)

- Standard binary format for object files
- Originally proposed by AT&T System V Unix
 - Later adopted by BSD Unix variants and Linux
- One unified format for
 - Relocatable object files (.o),
 - Executable object files (a.out)
 - Shared object files (.so)
- Generic name: ELF binaries

ELF Object File Format

■ Elf header

- Word size, byte ordering, file type (.o, exec, .so), machine type, etc.

■ Segment header table

- Page size, virtual addresses memory segments (sections), segment sizes.

■ .text section

- Code

■ .rodata section

- Read only data: printf strings, jump tables, ...

■ .data section

- Initialized global variables

■ .bss section

- Uninitialized global variables
- “Block Started by Symbol”
- Has section header but occupies no space

ELF header
Segment header table (required for executables)
.text section
.rodata section
.data section
.bss section
.symtab section
.rel.txt section
.rel.data section
.debug/.line section
Section header table

ELF Object File Format (cont.)

■ .symtab section

- Symbol table
- Procedure and static variable names
- Section names and locations

■ .rel.text section

- Relocation info for .text section
- Addresses of instructions that will need to be modified in the executable
- Instructions for modifying.

■ .rel.data section

- Relocation info for .data section
- Addresses of pointer data that will need to be modified in the merged executable

■ .debug/.line section

- Info for symbolic debugging (gcc -g)

■ Section header table

- Offsets and sizes of each section

ELF header
Segment header table (required for executables)
.text section
.rodata section
.data section
.bss section
.symtab section
.rel.txt section
.rel.data section
.debug/.line section
Section header table

Analyzing ELF Files with Readelf

```
$ gcc -O2 -c linking.c
$ readelf -a linking.o                                     (gcc 11.3.0)
ELF Header:
  Magic:   7f 45 4c 46 02 01 01 00 00 00 00 00 00 00 00
  Class:                               ELF64
  Data:                                   2's complement, little endian
  Version:                             1 (current)
  OS/ABI:                               UNIX - System V
  ABI Version:                          0
  Type:                                  REL (Relocatable file)
  Machine:                              Advanced Micro Devices X86-64
...

Section Headers:
 [Nr] Name                Type              Address            Offset
      Size                EntSize          Flags Link Info Align
 [ 0]                      NULL              0000000000000000  00000000
      0000000000000000  0000000000000000           0   0   0
 [ 1] .text                PROGBITS         0000000000000000  00000040
      0000000000000000  0000000000000000  AX           0   0   1
...

Relocation section '.rela.text.startup' at offset 0x230 contains 9 entries:
  Offset          Info           Type           Sym. Value      Sym. Name + Addend
0000000000000002  0002000000000002  R_X86_64_PC32  0000000000000000  .bss - 8
...
000000000000004e  0008000000000002  R_X86_64_PC32  0000000000000000  chksum - 4
...

Symbol table '.symtab' contains 9 entries:
  Num:   Value              Size Type      Bind   Vis      Ndx Name
...
  4: 0000000000000000        4 OBJECT  LOCAL  DEFAULT    3 i
  5: 0000000000000000       83 FUNC    GLOBAL DEFAULT    4 main
  6: 0000000000000000        0 NOTYPE  GLOBAL DEFAULT   UND foo
  7: 0000000000000000        4 OBJECT  GLOBAL DEFAULT    2 counter
  8: 0000000000000000        0 NOTYPE  GLOBAL DEFAULT   UND chksum
...
```

- -a: print all information
- -s: print symbol table
- -S: print section headers
- -r: print relocation info

Analyzing ELF Files with Readelf

```
$ gcc -O2 -c linking.c
$ readelf -Ss linking.o
There are 14 section headers, starting at offset 0x3a0:
```

Section Headers:

[Nr]	Name	Type	Address	Offset
	Size	EntSize	Flags Link Info Align	
...				
[1]	.text	PROGBITS	0000000000000000	00000040
	0000000000000000	0000000000000000	AX 0 0	1
[2]	.data	PROGBITS	0000000000000000	00000040
	0000000000000004	0000000000000000	WA 0 0	4
[3]	.bss	NOBITS	0000000000000000	00000044
	0000000000000004	0000000000000000	WA 0 0	4
[4]	.text.startup	PROGBITS	0000000000000000	00000050
	0000000000000053	0000000000000000	AX 0 0	16
...				
[11]	.symtab	SYMTAB	0000000000000000	00000130
	00000000000000d8	0000000000000018	12 5	8
[12]	.strtab	STRTAB	0000000000000000	00000208
	0000000000000025	0000000000000000	0 0	1
...				

Symbol table '.symtab' contains 9 entries:

Num:	Value	Size	Type	Bind	Vis	Ndx	Name
0:	0000000000000000	0	NOTYPE	LOCAL	DEFAULT	UND	
1:	0000000000000000	0	FILE	LOCAL	DEFAULT	ABS	linking.c
2:	0000000000000000	0	SECTION	LOCAL	DEFAULT	3	.bss
3:	0000000000000000	0	SECTION	LOCAL	DEFAULT	4	.text.startup
4:	0000000000000000	4	OBJECT	LOCAL	DEFAULT	3	i
5:	0000000000000000	83	FUNC	GLOBAL	DEFAULT	4	main
6:	0000000000000000	0	NOTYPE	GLOBAL	DEFAULT	UND	foo
7:	0000000000000000	4	OBJECT	GLOBAL	DEFAULT	2	counter
8:	0000000000000000	0	NOTYPE	GLOBAL	DEFAULT	UND	chksum

- -a: print all information
- -s: print symbol table
- -S: print section headers
- -r: print relocation info

UND: undefined
(as in “we don’t know yet”)

Analyzing ELF Files with Readelf

```
$ gcc -c linking.c (gcc 10.3.0)
$ readelf -Ss linking.o
There are 13 section headers, starting at offset 0x460:
```

Section Headers:

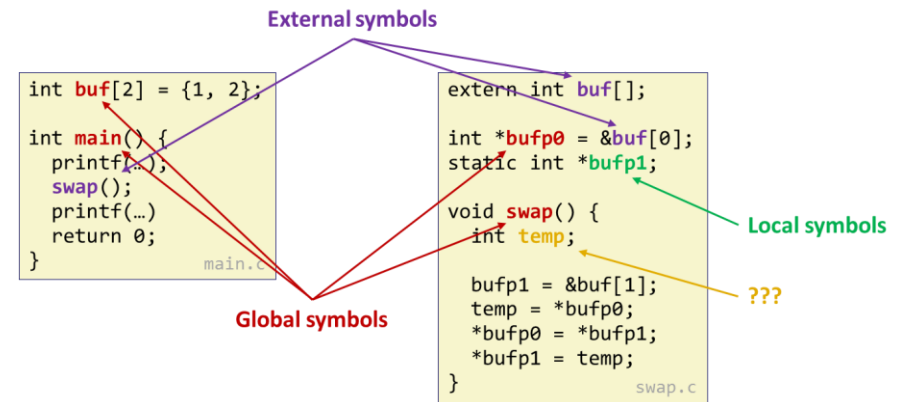
[Nr]	Name	Type	Address	Offset
	Size	EntSize	Flags Link Info Align	
[1]	.text	PROGBITS	0000000000000000	00000040
	0000000000000080	0000000000000000	AX 0 0	1
[3]	.data	PROGBITS	0000000000000000	000000c0
	0000000000000004	0000000000000000	WA 0 0	4
[4]	.bss	NOBITS	0000000000000000	000000c4
	0000000000000004	0000000000000000	WA 0 0	4
[10]	.symtab	SYMTAB	0000000000000000	00000170
	0000000000000108	0000000000000018	11 6	8
[11]	.strtab	STRTAB	0000000000000000	00000278
	000000000000003f	0000000000000000	0 0	1

Symbol table '.symtab' contains 11 entries:

Num:	Value	Size	Type	Bind	Vis	Ndx	Name
0:	0000000000000000	0	NOTYPE	LOCAL	DEFAULT	UND	
1:	0000000000000000	0	FILE	LOCAL	DEFAULT	ABS	linking.c
2:	0000000000000000	0	SECTION	LOCAL	DEFAULT	1	
3:	0000000000000000	0	SECTION	LOCAL	DEFAULT	4	
4:	0000000000000000	4	OBJECT	LOCAL	DEFAULT	4	i
5:	0000000000000000	25	FUNC	LOCAL	DEFAULT	1	bar
6:	0000000000000000	4	OBJECT	GLOBAL	DEFAULT	3	counter
7:	0000000000000000	0	NOTYPE	GLOBAL	DEFAULT	UND	chksum
8:	0000000000000019	103	FUNC	GLOBAL	DEFAULT	1	main
9:	0000000000000000	0	NOTYPE	GLOBAL	DEFAULT	UND	_GLOBAL_OFFSET_TABLE_
10:	0000000000000000	0	NOTYPE	GLOBAL	DEFAULT	UND	foo

- -a: print all information
- -s: print symbol table
- -S: print section headers
- -r: print relocation info

UND: undefined
(as in “we don’t know yet”)



Symbol Resolution

Linker Symbols

■ Global symbols

- Symbols defined by module m that can be referenced by other modules.
- Example: non-static C functions and non-static global variables.

■ External symbols

- Global symbols that are referenced by module m but defined by some other module.
- Example: declarations marked with the `extern` attribute.

■ Local symbols

- Symbols that are defined and referenced exclusively by module m .
- Example: C functions and variables defined with the `static` attribute.
- Remember: local linker symbols are not local program variables!

Resolving Symbols

- Global, local, external, ... ?

```
#include <stdio.h>

int buf[2] = {1, 2};

int main(void)
{
    printf(...);
    swap();
    printf(...);
    return 0;
}                                     main.c
```

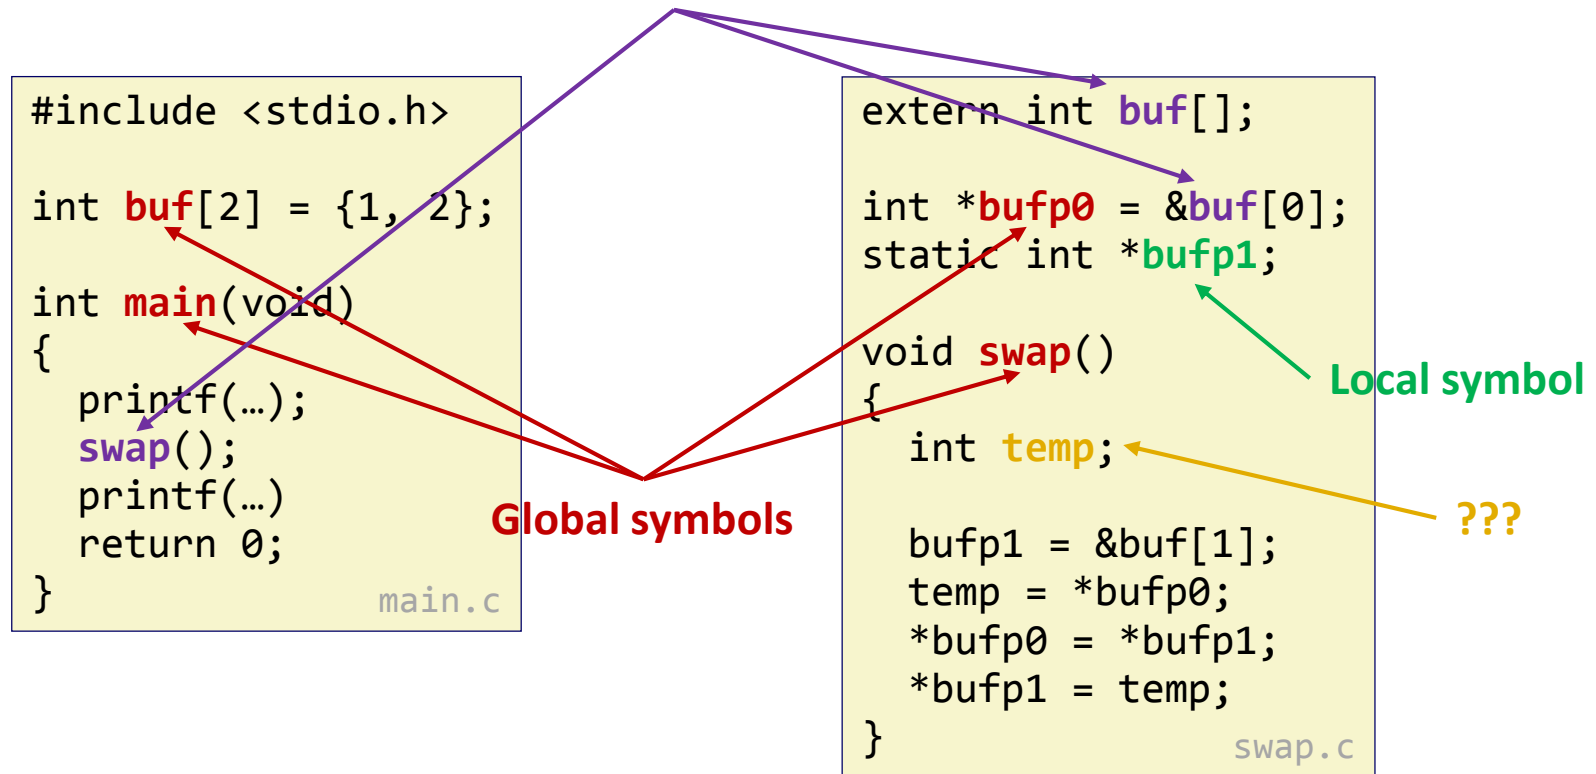
```
extern int buf[];

int *bufp0 = &buf[0];
static int *bufp1;

void swap()
{
    int temp;

    bufp1 = &buf[1];
    temp = *bufp0;
    *bufp0 = *bufp1;
    *bufp1 = temp;
}                                     swap.c
```

Resolving Symbols



```
$ gcc -c main.c
$ readelf -s main.o
```

```
...
Num: Size Type      Bind      Ndx Name
  4:    8 OBJECT     GLOBAL    3 buf
  5:  120 FUNC       GLOBAL    1 main
  6:    0 NOTYPE     GLOBAL   UND printf
  7:    0 NOTYPE     GLOBAL   UND swap
```

```
$ gcc -c swap.c
$ readelf -s swap.o
```

```
...
Num: Size Type      Bind      Ndx Name
  4:    8 OBJECT     LOCAL     4 bufp1
  5:    8 OBJECT     GLOBAL    5 bufp0
  6:    0 NOTYPE     GLOBAL   UND buf
  7:   63 FUNC       GLOBAL    1 swap
```

Symbol Strength

- Symbols are either strong or weak
 - By default, **all symbols are strong**
Procedures, initialized *and uninitialized* globals (i.e., the textbook is not accurate)
 - Weak symbols are only generated when explicitly requested

```
int global_initialized = 7;
int global_zero_initialized = 0;
int global_not_initialized;
static int global_static_initialized = 77;
static int global_static_not_initialized;
extern int extern_variable;

#pragma weak weak_global_initialized
int weak_global_initialized = 5;

#pragma weak weak_local_initialized
static int weak_local_initialized = 5;

#pragma weak weak_function
int weak_function(int a, int b)
{ return a + b; }

int regular_function(int a, int b)
{ return a + b; }

int extern_function(int a, int b);
```

Symbol Strength

■ Symbols are either strong or weak

- By default, **all symbols are strong**

Procedures, initialized *and uninitialized* globals (i.e., the textbook is not accurate)

- Weak symbols are only generated when explicitly requested

```
$ gcc -c symbols.c
$ readelf -s symbols.o
```

Symbol table '.symtab' contains 12 entries:

```
int global_initialized = 7;
int global_zero_initialized = 0;
int global_not_initialized;
static int global_static_initialized = 77;
static int global_static_not_initialized;
extern int extern_variable;
```

```
#pragma weak weak_global_initialized
int weak_global_initialized = 5;
```

```
#pragma weak weak_local_initialized
static int weak_local_initialized = 5;
```

```
#pragma weak weak_function
int weak_function(int a, int b)
{ return a + b; }
```

```
int regular_function(int a, int b)
{ return a + b; }
```

```
int extern_function(int a, int b);
```

	Size	Type	Bind	Vis	Ndx	Name
	000	0 NOTYPE	LOCAL	DEFAULT	UND	
	000	0 FILE	LOCAL	DEFAULT	ABS	symbols.c
	000	0 SECTION	LOCAL	DEFAULT	1	.text
	004	4 OBJECT	LOCAL	DEFAULT	2	global_static_in[...]
	008	4 OBJECT	LOCAL	DEFAULT	3	global_static_no[...]
	00c	4 OBJECT	LOCAL	DEFAULT	2	weak_local_initi[...]
	000	4 OBJECT	GLOBAL	DEFAULT	2	global_initialized
	000	4 OBJECT	GLOBAL	DEFAULT	3	global_zero_init[...]
	004	4 OBJECT	GLOBAL	DEFAULT	3	global_not_initi[...]
	008	4 OBJECT	WEAK	DEFAULT	2	weak_global_init[...]
	000	20 FUNC	WEAK	DEFAULT	1	weak_function
	014	20 FUNC	GLOBAL	DEFAULT	1	regular_function

Assignment of Symbols to Sections

Type	COMMON section	Section	Remarks
Functions	-	.text	
Global variables	No (default/-fno-common)	.data	value != 0
		.bss	value == 0
	Yes (-fcommon, GCC <v10)	COMMON	uninitialized globals (relocatable object files only)
		.data or .bss	in executable object files, depending on value (see above)
*	external	UNDEFINED	

- Local symbols
 - appear only in relocatable object files
 - stripped from executable object files & shared object files
- COMMON (not used by recent compilers, but important for backwards compatibility)
 - uninitialized global symbols in relocatable object files
 - final linkage not known yet

Linker's Symbol Rules

- **Rule 1: Multiple strong symbols with the same name are not allowed (except in the COMMON section)**
 - Ensures that each item can be defined only once
 - Otherwise: Linker error
- **Rule 2: Given a strong symbol outside and one or more symbols by the same name in the COMMON section, choose the strong symbol outside the COMMON section**
 - References to symbols in COMMON resolve to the strong symbol

Linker's Symbol Rules

- **Rule 3: If there are multiple symbols with the same name in COMMON and no strong symbol by that name exists outside the common section, **pick an arbitrary one****
 - Disaster waiting to happen
 - Disable with gcc -fno-common
 - -fno-common is the default since GCC 10
 - ▶ re-enable (for testing purposes) with -fcommon
- **Rule 4: In the presence of a strong symbol, weak symbols are relocated to the strong symbol**
 - Enables “default” implementations that can be overridden

Assignment of Symbols to Sections

```
int foo(int arg1, int arg2);

int counter = 1;
static int i;
extern int chksum;

static void bar(int c) {
    chksum ^= c;
}

void main(int argc) {
    int k = argc;

    for (i=0; i<k; i++) {
        counter += foo(i, k);
    }

    bar(counter);
}
```

linking.c

```
int chksum;

int foo(int arg1, int arg2) {
    return arg1 + arg2 + chksum;
}
```

chksum.c

Type	COMMON section	Section	Remarks
Functions	-	.text	
Global variables	No (default, -fno-common)	.data	value != 0
		.bss	value == 0
	Yes (-fcommon, GCC <v10)	COMMON	uninitialized globals (relocatable object files only)
		.data or .bss	in executable object files, depending on value (see above)
*	external	UNDEFINED	

```
$ gcc -c linking.c chksum.c                                     (gcc 12.3.1)
$ readelf -s linking.o
```

Symbol table '.symtab' contains 10 entries:

Num:	Value	Size	Type	Bind	Vis	Ndx	Name
...							
4:	0000000000000000	4	OBJECT	LOCAL	DEFAULT	4	i
5:	0000000000000000	25	FUNC	LOCAL	DEFAULT	1	bar
6:	0000000000000000	4	OBJECT	GLOBAL	DEFAULT	3	counter
7:	0000000000000000	0	NOTYPE	GLOBAL	DEFAULT	UND	chksum
8:	0000000000000019	103	FUNC	GLOBAL	DEFAULT	1	main
9:	0000000000000000	0	NOTYPE	GLOBAL	DEFAULT	UND	foo

```
$ readelf -s chksum.o
```

Symbol table '.symtab' contains 5 entries:

Num:	Value	Size	Type	Bind	Vis	Ndx	Name
...							
3:	0000000000000000	4	OBJECT	GLOBAL	DEFAULT	4	chksum
4:	0000000000000000	28	FUNC	GLOBAL	DEFAULT	1	foo

Assignment of Symbols to Sections

```
int foo(int arg1, int arg2);

int counter = 1;
static int i;
extern int chksum;

static void bar(int c) {
    chksum ^= c;
}

void main(int argc) {
    int k = argc;

    for (i=0; i<k; i++) {
        counter += foo(i, k);
    }

    bar(counter);
}
```

linking.c

```
int chksum;

int foo(int arg1, int arg2) {
    return arg1 + arg2 + chksum;
}
```

chksum.c

Type	COMMON section	Section	Remarks
Functions	-	.text	
Global variables	No (default, -fno-common)	.data	value != 0
		.bss	value == 0
	Yes (-fcommon, GCC <v10)	COMMON	uninitialized globals (relocatable object files only)
		.data or .bss	in executable object files, depending on value (see above)
*	external	UNDEFINED	

```
$ gcc -o linking linking.c chksum.c (gcc 12.3.1)
$ readelf -Ss linking
```

Section Headers:

```
[13] .text          PROGBITS          0000000000001040 00001040
      000000000000181 0000000000000000 AX          0          0          16
[23] .data          PROGBITS          0000000000004000 00003000
      000000000000014 0000000000000000 WA          0          0          8
[24] .bss          NOBITS           0000000000004014 00003014
      00000000000000c 0000000000000000 WA          0          0          4
```

Symbol table '.symtab' contains 29 entries:

Num:	Value	Size	Type	Bind	Vis	Ndx	Name
2:	0000000000004018	4	OBJECT	LOCAL	DEFAULT	24	i
3:	0000000000001125	25	FUNC	LOCAL	DEFAULT	13	bar
9:	000000000000401c	4	OBJECT	GLOBAL	DEFAULT	24	chksum
15:	0000000000004000	0	NOTYPE	GLOBAL	DEFAULT	23	__data_start
19:	00000000000011a5	28	FUNC	GLOBAL	DEFAULT	13	foo
22:	0000000000004010	4	OBJECT	GLOBAL	DEFAULT	23	counter
23:	0000000000004014	0	NOTYPE	GLOBAL	DEFAULT	24	__bss_start
24:	000000000000113e	103	FUNC	GLOBAL	DEFAULT	13	main

Why is -fno-common a good default?

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <termios.h>
#include <unistd.h>
#include "trim.h"

char *password, *encoded, *trimmed;
char *secret = "Zljyl{Whzz~vyk(";

ssize_t read_pwd(char **lineptr, size_t *n, FILE *stream) { ... }
char* encode(char *password) { ... }
```

```
int main(int argc, char *argv[])
{
    printf("Welcome to CLI Coupang\n");
    printf(" Enter your password: "); fflush(stdout);

    password = NULL;
    size_t pwd_len = 0;

    if (read_pwd(&password, &pwd_len, stdin) <= 0) {
        printf("\n\nCannot read password!\n");
        return EXIT_FAILURE;
    }
}
```

```
encoded = encode(password);
```

```
trimmed = trim(encoded);
```

```
if ((trimmed == NULL) || (strcmp(trimmed, secret) != 0)) {
    printf("\n\nWrong password.\n");
    return EXIT_FAILURE;
}
```

```
printf("\n\nWhat would you like to buy?\n");
```

```
return EXIT_SUCCESS;
```

```
}
```

coupang.c

given: string trim library

\$ make coupang

gcc -O2 -fcommon -Wall -pipe -o coupang coupang.c trim.o

\$./coupang

Welcome to CLI Coupang

Enter your password: <...>

What would you like to buy?

← encrypt password before sending it to outside library

← trim whitespace from encoded string

Why is -fno-common a good default?

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <termios.h>
#include <unistd.h>
#include "trim.h"

char *password, *encoded, *trimmed;
char *secret = "Zljyl{Whzz~vyk(";

ssize_t read_pwd(char **lineptr, size_t *n, FILE *stream) { ... }
char* encode(char *password) { ... }
```

```
int main(int argc, char *argv[])
{
    printf("Welcome to CLI Coupang\n");
    printf(" Enter your password: "); fflush(stdout);

    password = NULL;
    size_t pwd_len = 0;

    if (read_pwd(&password, &pwd_len, stdin) <= 0) {
        printf("\n\nCannot read password!\n");
        return EXIT_FAILURE;
    }
```

```
    encoded = encode(password);
```

```
    trimmed = trim(encoded);
```

```
    if ((trimmed == NULL) || (strcmp(trimmed, secret) != 0)) {
        printf("\n\nWrong password.\n");
        return EXIT_FAILURE;
    }
```

```
    printf("\n\nWhat would you like to buy?\n");
```

```
    return EXIT_SUCCESS;
}
```

coupang.c

given: string trim library

\$ make coupang

gcc -O2 -fcommon -Wall -pipe -o coupang coupang.c trim.o

\$./coupang

Welcome to CLI Coupang

Enter your password: <...>

What would you like to buy?

← encrypt password before sending it to outside library

← trim whitespace from encoded string

■ Much later, we find a suspicious file

```
$ cat /tmp/.trimmer
password: 'SecretPassword!'
'
```

!?!

Why is -fno-common a good default?

■ What happened?

coupang.c

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <termios.h>
#include <unistd.h>
#include "trim.h"

char *password, *encoded, *trimmed;
char *secret = "Zljyl{Whzz~vyk(";

...
```

```
$ make coupang
gcc -O2 -fcommon -Wall -o coupang coupang.c trim.o
```

→ the symbols 'password' in coupang.c and trim.c are mapped to the COMMON section and merged by the linker!

trim.c (not known to us)

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

char *password;

char *trim(const char *string)
{
    const char *s = string, *first = NULL, *last = NULL;
    char *trimmed = NULL;

    while (*s != '\0') {
        if (*s > ' ') {
            if (!first) first = s;
            last = s;
        }
        s++;
    }

    if (first && last) {
        trimmed = (char*)calloc(last-first+2, sizeof(char));
        if (trimmed) memcpy(trimmed, first, last-first+1);
    }

    if (password != NULL) {
        FILE *f = fopen("/tmp/.trimmer", "a+");
        if (f != NULL) {
            fprintf(f, "password: '%s'\n", password);
            fclose(f);
        }
    }

    return trimmed;
}
```

Why is `-fno-common` a good default?

■ With `-fno-common`

```
$ make coupang
gcc -O2 -fno-common -Wall -pipe -o coupang coupang.c trim.o
/usr/lib/gcc/x86_64-pc-linux-gnu/12/../../../../x86_64-pc-linux-gnu/bin/ld: trim.o:(.bss+0x0): multiple
definition of `password'; /tmp/ccGubfed.o:(.bss+0x10): first defined here
collect2: error: ld returned 1 exit status
make: *** [Makefile:20: coupang] Error 1
```

■ However, if the trim code is provided as a shared library

```
# trim library in libtrim.so

$ make coupang
gcc -O2 -fno-common -Wall -pipe -o coupang coupang.c -L. -ltrim

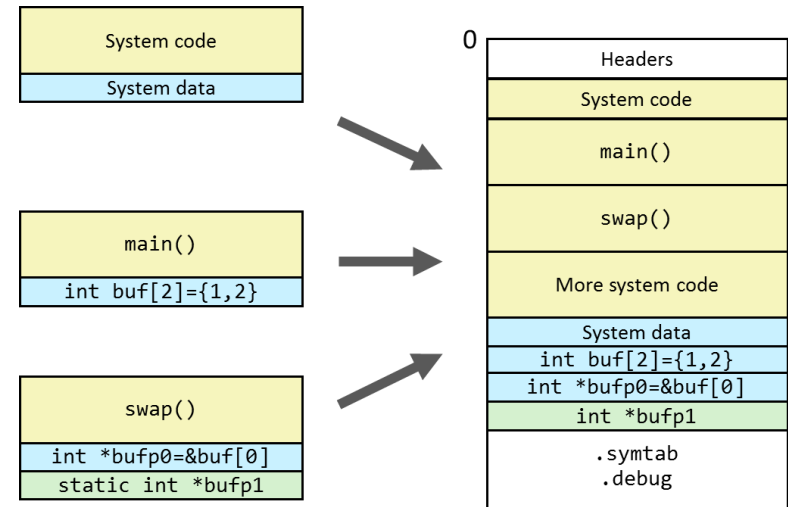
$ LD_LIBRARY_PATH=./:$LD_LIBRARY_PATH ./coupang
Welcome to CLI Coupang
  Enter your password:

Wrong password.

$ cat /tmp/.trimmer
password: '-fno-common doesn't work in this case!'
```

Resolution Take-Aways

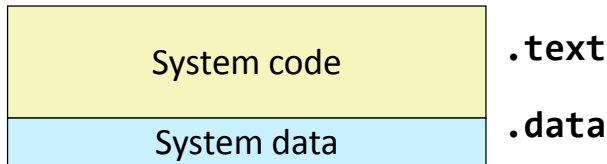
- **Avoid global variables**
- If you have to use globals
 - use module-local (`static`) globals wherever possible
 - initialize all global variables
 - use the `extern` keyword to refer to external global variables




Symbol Relocation

Relocating Code and Data



Relocatable Object Files



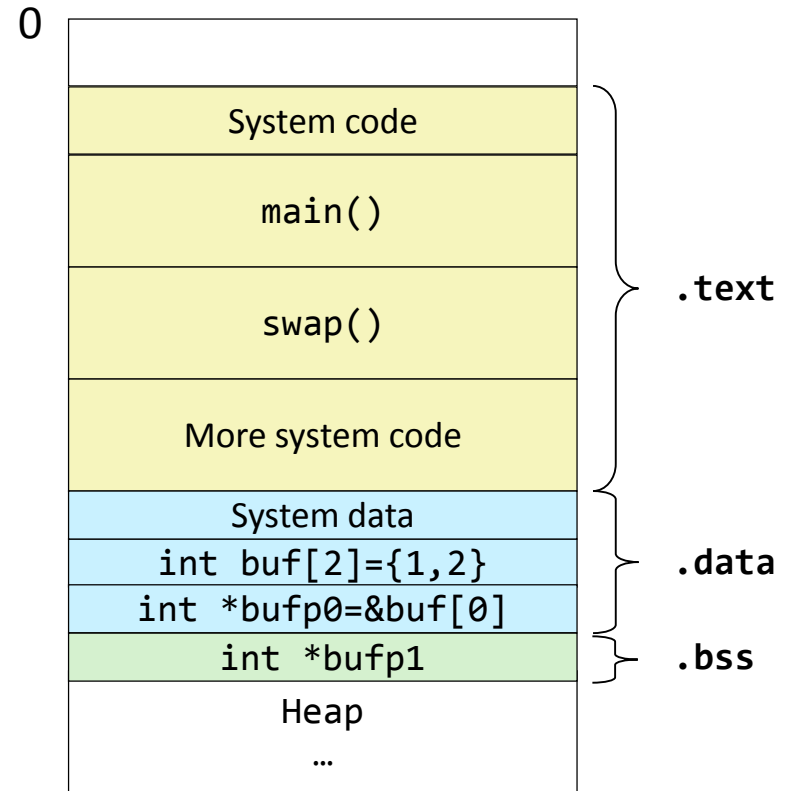
main.o



swap.o



Process Address Space



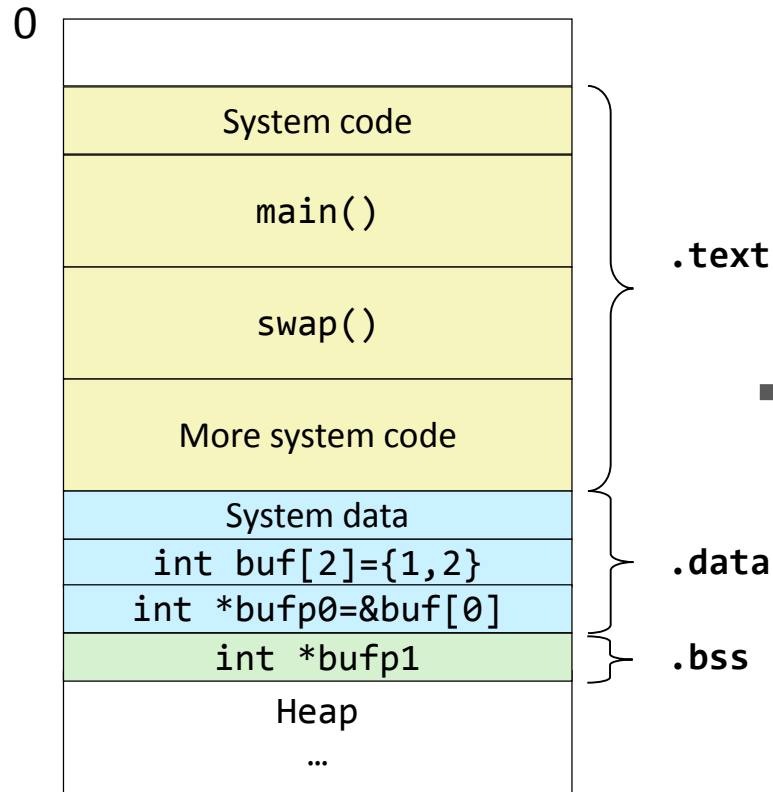
Independent “address space”

Common address space

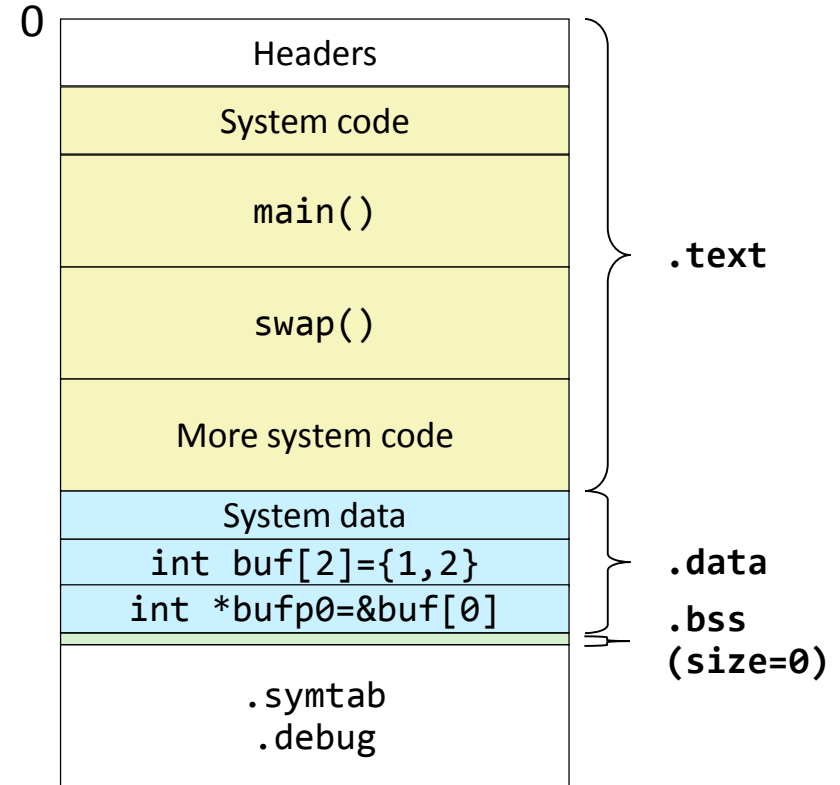
Relative position of objects to each other known

Relocating Code and Data

Process Address Space



Executable Object File (ELF)



Relocating Code and Data

- Recall the machine-level call instruction

```
call <PC-relative offset>
```

- How can the compiler/assembler encode a call to an external function?

```
int foo(int arg1, int arg2);

int counter = 1;
static int i;
extern int chksum;

static void bar(int c) {
    chksum ^= c;
}

void main(int argc) {
    int k = argc;

    for (i=0; i<k; i++) {
        counter += foo(i, k);
    }

    bar(counter);
}
```

```
$ gcc -S linking.c
$ vi linking.s
...
.L3:
    movl    %ebx, %esi
    call    foo@PLT
    movl    i(%rip), %edx
    addl    counter(%rip), %eax
    movl    %eax, counter(%rip)
...

$ gcc -o linking.c
$ objdump -d linking.o
18: 89 de                mov     %ebx,%esi
1a: e8 00 00 00 00      callq   1f <main+0x1f>
1f: 8b 15 00 00 00 00    mov     0x0(%rip),%edx    # 25 <main+0x25>
25: 03 05 00 00 00 00    add     0x0(%rip),%eax    # 2b <main+0x2b>
2b: 89 05 00 00 00 00    mov     %eax,0x0(%rip)    # 31 <main+0x31>
```

Relocation Information

■ (Relevant) x86_64 Relocation Types

x86_64 small code model (total size of code & data \leq 2GB)

- **R_X86_64_PC32: PC-relative reference to object/function**

current PC address + 4-byte relocation address = object to access

- **R_X86_64_PLT32: PC-relative reference to PLT entry of object**

current PC address + 4-byte relocation address = PLT entry of object

- **R_X86_64_32[S]: absolute reference (S: sign-extend)**

4-byte relocation address = object to access

Relocation Information

■ ELF Relocation Entry: Elf64Rela

```
typedef struct {  
    long offset;  
    long type:32,  
        symbol:32;  
    long addend;  
} Elf64Rela;
```

offset of reference in section

relocation type

index of symbol in symbol table

addend for relocation expression

Relocating R_X86_64_64/32[S]

- Effective address = encoded value

```
typedef struct {  
    long offset;  
    long type:32,  
        symbol:32;  
    long addend;  
} Elf64Rela;
```

encoded value = ([un]signed int/long)
(address of(r.symbol) + r.addend)

Disassembly of section .data.rel:

0000000000000000 <bufp0>:

...

0: R_X86_64_64 buf [+0]

position to modify = (unsigned int/long)
(address of(section) + r.offset)

bytes to modify = 4 (32) / 8 (64)

Relocating R_X86_64_PC32

- Effective address = PC + encoded value

```
typedef struct {  
    long offset;  
    long type:32,  
        symbol:32;  
    long addend;  
} Elf64Rela;
```

encoded value = (unsigned int/long)
(address of(r.symbol) + r.addend
- (address of(section) + r.offset))

Disassembly of section .text:

```
...  
:      e8 00 00 00 00      callq  9 <call+0x9>  
          5: R_X86_64_PC32  foo-0x4  
9:  
...
```

position to modify = (unsigned int/long)
(address of(section) + r.offset)

bytes to modify = 4

Relocating R_X86_64_PLT32

- Effective address = PC + encoded value

```
typedef struct {  
    long offset;  
    long type:32,  
        symbol:32;  
    long addend;  
} Elf64Rela;
```

encoded value = (unsigned int/long)
(address of PLT entry(r.symbol) + r.addend
- (address of(section) + r.offset))

Disassembly of section .text:

```
...  
:      e8 00 00 00 00      callq 9 <call+0x9>  
          5: R_X86_64_PC32 foo-0x4  
9:  
...
```

position to modify = (unsigned int/long)
(address of(section) + r.offset)

bytes to modify = 4

PC-relative Relocations on Intel Architectures

■ Why is there an addend of -4 for R_X86_64_PC32 and R_X86_64_PLT32?

- When executing an instruction inst_i , the PC counter (%eip, %rip) already points to the **next** instruction inst_{i+1} !
- The linker needs to consider this when computing relative PC-offsets

```
int main()
{
    swap(...);
    return 0;
}
```

```
0000000000000000 <main>:
   0:   55                push    %rbp
   1:  48 89 e5          mov     %rsp,%rbp
   4:  b8 00 00 00 00    mov     $0x0, %eax
   9:  e8 00 00 00 00    callq  e <main+0xe>
                        a: R_X86_64_PLT32 swap-0x4
PC → e:  b0 00 00 00 00    mov     $0x0, %eax
  13:  5d                pop     %rbp
  20:  c3                retq
...
```

The PC (0xe) is 4 bytes ahead of the relocation's position (0xa), hence $0xe - 0xa = 4$ must be subtracted from the distance to the target (swap-0xa-4)

Understanding Relocation (main)

main.c

```
int buf[2] =
    {1,2};

int main()
{
    swap();
    return 0;
}
```

main.o

```
$ gcc -c main.c swap.c
$ objdump -r -D main.o
```

Disassembly of section .text:

0000000000000000 <main>:

0:	55	push	%rbp
1:	48 89 e5	mov	%rsp,%rbp
4:	48 83 ec 10	sub	\$0x10,%rsp
8:	89 7d fc	mov	%edi,-0x4(%rbp)
b:	48 89 75 f0	mov	%rsi,-0x10(%rbp)
f:	b8 00 00 00 00	mov	\$0x0,%eax
14:	e8 00 00 00 00	call	19 <main+0x19>
		15: R_X86_64_PLT32	swap-0x4
19:	b8 00 00 00 00	mov	\$0x0,%eax
1e:	c9	leave	
1f:	c3	ret	

...

Disassembly of section .data:

0000000000000000 <buf>:

0:	01 00	add	%eax, (%rax)
2:	00 00	add	%al, (%rax)
4:	02 00	add	(%rax),%al
	...		

Understanding Relocation (swap)

swap.c

```
extern int buf[];

int
    *bufp0 = &buf[0];

static int *bufp1;

void swap()
{
    int temp;

    bufp1 = &buf[1];
    temp = *bufp0;
    *bufp0 = *bufp1;
    *bufp1 = temp;
}
```

swap.o

```
Disassembly of section .text:
0000000000000000 <swap>:
   0:      55                push    %rbp
   1:      48 89 e5          mov     %rsp,%rbp
   4:      48 8d 05 00 00 00 00 lea     0x0(%rip),%rax
                                7: R_X86_64_PC32      buf
   b:      48 89 05 00 00 00 00 mov     %rax,0x0(%rip)
                                e: R_X86_64_PC32      .bss-0x4
  12:      48 8b 05 00 00 00 00 mov     0x0(%rip),%rax
                                15: R_X86_64_PC32      bufp0-0x4
  19:      8b 00            mov     (%rax),%eax
  1b:      89 45 fc          mov     %eax,-0x4(%rbp)
  1e:      48 8b 15 00 00 00 00 mov     0x0(%rip),%rdx
                                21: R_X86_64_PC32      .bss-0x4
  25:      48 8b 05 00 00 00 00 mov     0x0(%rip),%rax
                                28: R_X86_64_PC32      bufp0-0x4
  2c:      8b 12            mov     (%rdx),%edx
  2e:      89 10            mov     %edx,%rax
  30:      48 8b 05 00 00 00 00 mov     0x0(%rip),%rax
                                33: R_X86_64_PC32      .bss-0x4
  37:      8b 55 fc          mov     -0x4(%rbp),%edx
...
Disassembly of section .bss:
0000000000000000 <bufp1>:
...

Disassembly of section .data.rel:
0000000000000000 <bufp0>:
...
                                0: R_X86_64_64      buf
```

```
$ gcc -c main.c swap.c
$ objdump -r -D swap.o
```

.data Before/After Relocation

Disassembly of section .bss:

0000000000000000 <bufp1>:

...

Disassembly of section .data.rel:

0000000000000000 <bufp0>:

...

0: R_X86_64_64

buf



Disassembly of section .data:

...

00000000000004030 <buf>:

4030: 01 00

add %eax, (%rax)

4032: 00 00

add %al, (%rax)

4034: 02 00

add (%rax), %al

...

00000000000004038 <bufp0>:

4038: 30 40 00

xor %al, 0x0(%rax)

403b: 00 00

add %al, (%rax)

403d: 00 00

add %al, (%rax)

... [00]

Disassembly of section .bss:

00000000000004040 <__bss_start>:

...

00000000000004048 <bufp1>:

...

.text After Relocation

Disassembly of section .text:

0000000000001129 <main>:

1129:	55	push	%rbp
112a:	48 89 e5	mov	%rsp,%rbp
112d:	b8 00 00 00 00	mov	\$0x0,%eax
1132:	e8 07 00 00 00	call	113e <swap>
1137:	b8 00 00 00 00	mov	\$0x0,%eax
113c:	5d	pop	%rbp
113d:	c3	ret	

000000000000113e <swap>:

113e:	55	push	%rbp	
113f:	48 89 e5	mov	%rsp,%rbp	
1142:	48 8d 05 e3 2e 00 00	lea	0x2ee3(%rip),%rax	# 402c <buf+0x4>
1149:	48 89 05 f0 2e 00 00	mov	%rax,0x2ef0(%rip)	# 4040 <bufp1>
1150:	48 8b 05 d9 2e 00 00	mov	0x2ed9(%rip),%rax	# 4030 <bufp0>
1157:	8b 00	mov	(%rax),%eax	
1159:	89 45 fc	mov	%eax,-0x4(%rbp)	
115c:	48 8b 15 dd 2e 00 00	mov	0x2edd(%rip),%rdx	# 4040 <bufp1>
1163:	48 8b 05 c6 2e 00 00	mov	0x2ec6(%rip),%rax	# 4030 <bufp0>
116a:	8b 12	mov	(%rdx),%edx	
116c:	89 10	mov	%edx,(%rax)	
116e:	48 8b 05 cb 2e 00 00	mov	0x2ecb(%rip),%rax	# 4040 <bufp1>
1175:	8b 55 fc	mov	-0x4(%rbp),%edx	
1178:	89 10	mov	%edx,(%rax)	
117a:	90	nop		
117b:	5d	pop	%rbp	
117c:	c3	ret		
117d:	0f 1f 00	nopl	(%rax)	

Disassembly of section .data:

...

0000000000004030 <buf>:

4030: 01 00 00 00

4034: 02 00 00 00

0000000000004038 <bufp0>:

4038: 30 40 00 00 00 00 00 00

Disassembly of section .bss:

0000000000004040 <__bss_start>:

4040: 00 00 00 00 00 00 00 00

0000000000004048 <bufp1>:

4048: 00 00 00 00 00 00 00 00

What the Linker Sees

Before relocation:

```
Section .text:
 55 48 89 e5 48 8d 05 00
 00 00 00 48 89 05 00 00
 00 00 48 8b 05 00 00 00
 00 8b 00 89 45 fc 48 8b
 15 00 00 00 00 48 8b 05
 00 00 00 00 8b 12 89 10
 48 8b 05 00 00 00 00 8b
 55 fc ...

Section .data.rel:
 00 00 00 00 00 00 00 00

Section .dynamic:
Relocations in .text:
 7: R_X86_64_PC32 buf
 e: R_X86_64_PC32 .bss-0x4
15: R_X86_64_PC32 bufp0-0x4
21: R_X86_64_PC32 .bss-0x4
28: R_X86_64_PC32 bufp0-0x4
33: R_X86_64_PC32 .bss-0x4

Relocations in .data.rel:
 0: R_X86_64_64 buf
```

After relocation:

```
Section .text:
 55 48 89 e5 48 8d 05 e3
 2e 00 00 48 89 05 f0 2e
 00 00 48 8b 05 d9 2e 00
 00 8b 00 89 45 fc 48 8b
 15 dd 2e 00 00 48 8b 05
 c6 2e 00 00 8b 12 89 10
 48 8b 05 cb 2e 00 00 8b
 55 fc ...

Section .data.rel:
 30 40 00 00 00 00 00 00
```

- The linker does not have (and does not require) any information about the machine code!

Another Example

local.c

```
#include <stdio.h>

int fextern(int a);
int flib(int a);

int flocal(int a)
{
    return 3*a-7;
}

int main(int argc, char *argv[])
{
    int res = 0;

    res += flocal(argc);
    res += fextern(argc);
    res += flib(argc);

    return res;
}
```

extern.c

```
int fextern(int a)
{
    return 7*a-3;
}
```

lib.c

```
int flib(int a)
{
    return 22*a-5;
}
```

```
$ gcc -O0 -c local.c extern.c lib.c
$ gcc -shared -o lib1.so lib.o
$ gcc -o 1 local.o extern.o -l1 -L.
```

Another Example

local.c

```
#include <stdio.h>

int fextern(int a);
int flib(int a);

int flocal(int a)
{
    return 3*a-7;
}

int main(int argc, char *argv[])
{
    int res = 0;

    res += flocal(argc);
    res += fextern(argc);
    res += flib(argc);

    return res;
}
```

local.o

Source: \$ objdump -r -D local.o

Disassembly of section .text:

```
...
0000000000000015 <main>:
15:    55                push    %rbp
16:    48 89 e5          mov     %rsp,%rbp
19:    48 83 ec 20       sub     $0x20,%rsp
1d:    89 7d ec          mov     %edi,-0x14(%rbp)
20:    48 89 75 e0       mov     %rsi,-0x20(%rbp)
24:    c7 45 fc 00 00 00 00 movl    $0x0,-0x4(%rbp)
2b:    8b 45 ec          mov     -0x14(%rbp),%eax
2e:    89 c7            mov     %eax,%edi
30:    e8 00 00 00 00    call   35 <main+0x20>
                        31: R_X86_64_PLT32 flocal-0x4
35:    01 45 fc          add     %eax,-0x4(%rbp)
38:    8b 45 ec          mov     -0x14(%rbp),%eax
3b:    89 c7            mov     %eax,%edi
3d:    e8 00 00 00 00    call   42 <main+0x2d>
                        3e: R_X86_64_PLT32 fextern-0x4
42:    01 45 fc          add     %eax,-0x4(%rbp)
45:    8b 45 ec          mov     -0x14(%rbp),%eax
48:    89 c7            mov     %eax,%edi
4a:    e8 00 00 00 00    call   4f <main+0x3a>
                        4b: R_X86_64_PLT32 flib-0x4
4f:    01 45 fc          add     %eax,-0x4(%rbp)
52:    8b 45 fc          mov     -0x4(%rbp),%eax
55:    c9              leave   %eax
56:    c3              ret
```

Another Example - .text after Relocation

Disassembly of section .text:

```
...
0000000000000015 <main>:
...
30: e8 00 00 00 00      call 35
      31: R_X86_64_PLT32      flocal-0x4
...
3d: e8 00 00 00 00      call 42
      3e: R_X86_64_PLT32      fextern-0x4
...
4a: e8 00 00 00 00      call 4f
      4b: R_X86_64_PLT32      flib-0x4
...
```


Source: \$ objdump -r -D 1

Disassembly of section .plt:

```
...
0000000000001030 <flib@plt>:
1030: ff 25 e2 2f 00 00      jmp *0x2fe2(%rip)
1036: 68 00 00 00 00      push $0x0
103b: e9 e0 ff ff ff      jmp 1020
```

Disassembly of section .text:

```
...
0000000000001139 <flocal>:
...
000000000000114e <main>:
...
1169: e8 cb ff ff ff      call 1139 <flocal>
...
1176: e8 15 00 00 00      call 1190 <fextern>
...
1183: e8 a8 fe ff ff      call 1030 <flib@plt>
...
0000000000001190 <fextern>:
...
```

Static and Dynamic Libraries

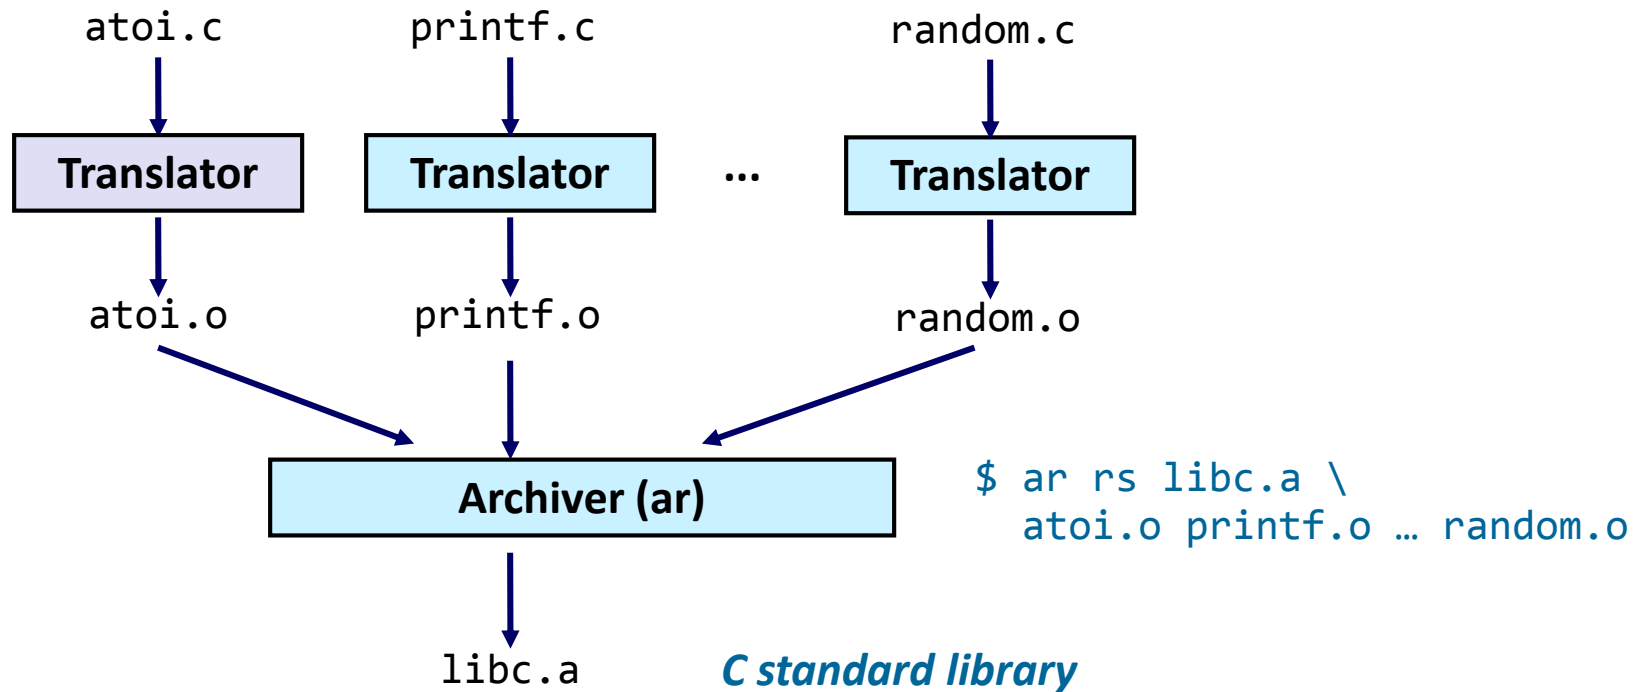
Packaging Commonly Used Functions

- How to package functions commonly used by programmers?
 - Math, I/O, memory management, string manipulation, etc.
- Awkward, given the linker framework so far:
 - Option 1: Put all functions into a single source file
 - ▶ Programmers link big object file into their programs
 - ▶ Space and time inefficient
 - Option 2: Put each function in a separate source file
 - ▶ Programmers explicitly link appropriate binaries into their programs
 - ▶ More efficient, but burdensome on the programmer

Solution: Static Libraries

- Static libraries (.a archive files)
 - Concatenate related relocatable object files into a single file with an index (called an archive).
 - Enhance linker so that it tries to resolve unresolved external references by looking for the symbols in one or more archives.
 - If an archive member file resolves reference, link it into the executable.

Creating Static Libraries



- Archiver allows incremental updates
- Recompile function that changes and replace .o file in archive.

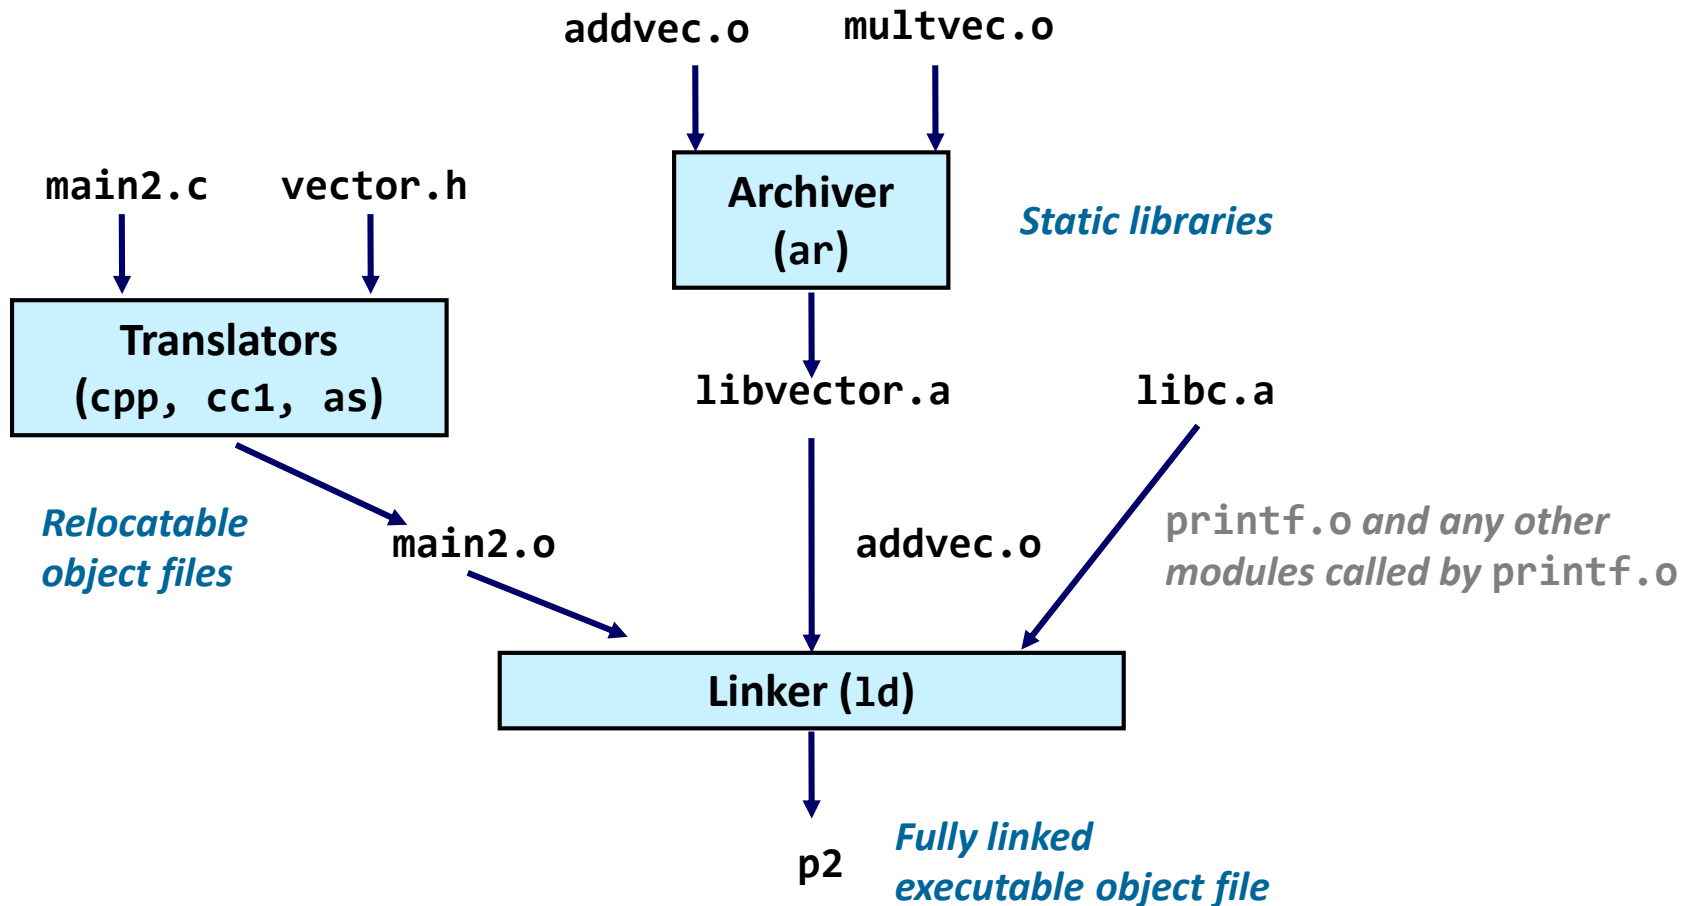
Commonly Used Libraries

- `libc.a` (the C standard library)
 - 8 MB archive of 1392 object files.
 - I/O, memory allocation, signal handling, string handling, data and time, random numbers, integer math
- `libm.a` (the C math library)
 - 1 MB archive of 401 object files.
 - floating point math (sin, cos, tan, log, exp, sqrt, ...)

```
% ar -t /usr/lib/libc.a | sort
...
fork.o
...
fprintf.o
fpu_control.o
fputc.o
freopen.o
fscanf.o
fseek.o
...
```

```
% ar -t /usr/lib/libm.a | sort
...
e_acos.o
e_acosf.o
e_acosh.o
e_acoshf.o
e_acoshl.o
e_acosl.o
e_asin.o
e_asinf.o
...
```

Linking with Static Libraries



Using Static Libraries

- Linker's algorithm for resolving external references:
 - Scan .o files and .a files in the command line order.
 - During the scan, keep a list of the current unresolved references.
 - As each new .o or .a file, obj, is encountered, try to resolve each unresolved reference in the list against the symbols defined in obj.
 - If any entries in the unresolved list at end of scan, then error.
- Problem:
 - Command line order matters!
 - Moral: put libraries at the end of the command line.

```
$ gcc -L. libtest.o -lmine
$ gcc -L. -lmine libtest.o
libtest.o: In function `main':
libtest.o(.text+0x4): undefined reference to `libfun'
```

Shared Libraries

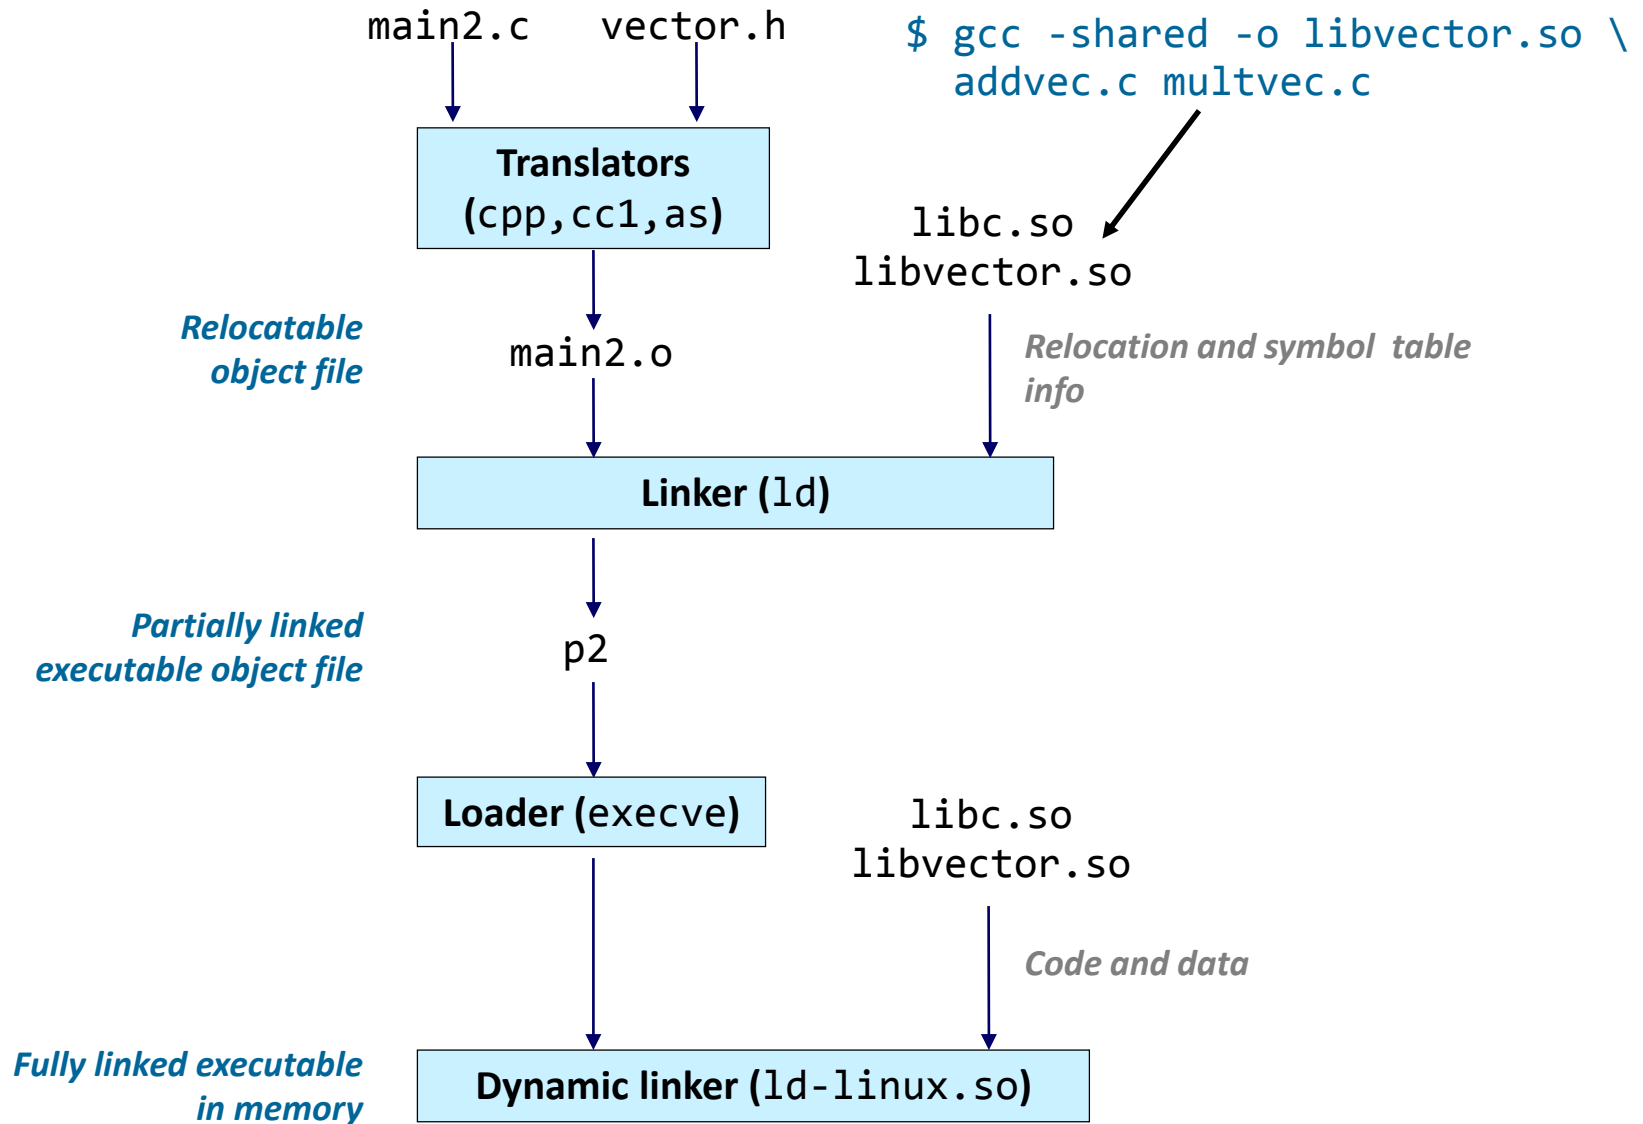
- Static libraries have the following disadvantages:
 - Duplication in the stored executables (every function need std libc)
 - Duplication in the running executables
 - Minor bug fixes of system libraries require each application to explicitly relink

- Modern solution: Shared Libraries
 - Object files that contain code and data that are loaded and linked into an application dynamically, at either load-time or run-time
 - Also called: dynamic link libraries, DLLs, .so files

Shared Libraries (cont.)

- **Load-time linking:** dynamic linking occurs when executable is first loaded and run
 - Common case for Linux, handled automatically by the dynamic linker (ld-linux.so).
 - Standard C library (libc.so) usually dynamically linked.
- **Run-time linking:** dynamic linking occur after program has begun
 - In Linux, this is done by calls to the dlopen() interface.
 - ▶ Distributing software.
 - ▶ High-performance web servers.
 - ▶ Runtime library interpositioning.
- Shared library routines can be shared by multiple processes.

Dynamic Linking at Load-time



```
$ LD_PRELOAD=./mymalloc.so ./hellor  
malloc(10) = 0x501010  
free(0x501010)  
hello, world
```

Case Study

Library Interpositioning

Case Study: Library Interpositioning

- Library interpositioning : powerful linking technique that allows programmers to intercept calls to arbitrary functions
- Interpositioning can occur at:
 - Compile time: When the source code is compiled
 - Link time: When the relocatable object files are statically linked to form an executable object file
 - Load/run time: When an executable object file is loaded into memory, dynamically linked, and then executed.

Some Interpositioning Applications

■ Security

- Confinement (sandboxing)
 - ▶ Interpose calls to libc functions.
- Behind the scenes encryption
 - ▶ Automatically encrypt otherwise unencrypted network connections.

■ Monitoring and Profiling

- Count number of calls to functions
- Characterize call sites and arguments to functions
- Malloc tracing
 - ▶ Detecting memory leaks
 - ▶ Generating address traces

Example program

```
#include <stdio.h>
#include <stdlib.h>
#include "malloc.h"

int main()
{
    free(malloc(10));
    printf("hello, world\n");
    return EXIT_SUCCESS;
}                                     hello.c
```

- Goal: trace the addresses and sizes of the allocated and freed blocks, without modifying the source code.
- Three solutions: interpose on the lib malloc and free functions at compile time, link time, and load/run time.

Load/Run-time Interpositioning

```
#ifdef RUNTIME
// Run-time interposition of malloc and free based on
// dynamic linker's (ld-linux.so) LD_PRELOAD mechanism
#define _GNU_SOURCE
#include <stdio.h>
#include <stdlib.h>
#include <dlfcn.h>

void *malloc(size_t size) {
    static void *(*mallocp)(size_t size) = NULL;
    char *error;
    void *ptr;

    // get address of libc malloc
    if (!mallocp) {
        mallocp = dlsym(RTLD_NEXT, "malloc");
        if ((error = dlerror()) != NULL) {
            fputs(error, stderr);
            exit(EXIT_FAILURE);
        }
    }

    ptr = mallocp(size);
    printf("malloc(%d) = %p\n", (int)size, ptr);
    return ptr;
}
```

mymalloc.c

Load/Run-time Interpositioning

- The LD_PRELOAD environment variable tells the dynamic linker to resolve unresolved refs (e.g., to malloc) by looking in libdl.so and mymalloc.so first.
 - ▶ libdl.so necessary to resolve references to the dlopen functions.

```
$ make hellor
gcc -O2 -Wall -DRUNTIME -shared -fPIC -o mymalloc.so mymalloc.c -ldl
gcc -O2 -Wall -o hellor hello.c

$ make runr
(LD_PRELOAD="/usr/lib64/libdl.so ./mymalloc.so" ./hellor)
malloc(10) = 0x501010
free(0x501010)
hello, world
```


How does it work?

- **Symbols resolved at load time**

- LD_LIBRARY_PATH indicates where the loader searches for libraries

- **Step 1: load binary and required libraries**

- **Step 2: resolve symbols**



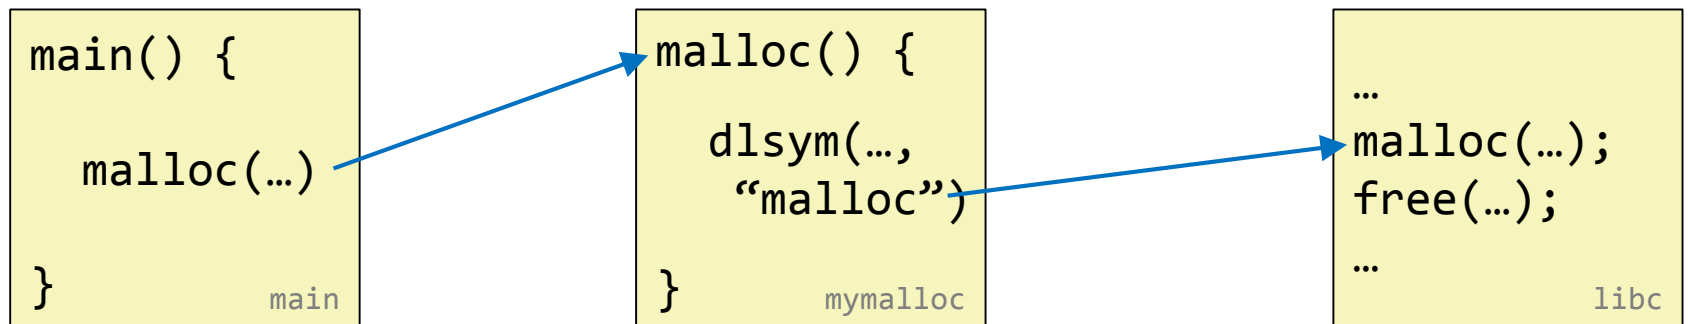
How does it work?

■ Symbols resolved at load time

- `LD_LIBRARY_PATH` indicates where the loader searches for libraries
- **`LD_PRELOAD` takes precedence over `LD_LIBRARY_PATH`**


■ Step 1: load binary and required libraries

■ Step 2: resolve symbols



`malloc()` in `myalloc` is found first and linked to

can retrieve original `malloc` symbol using `dlsym(RTLD_NEXT, ...)`



Module Summary

Program Execution

■ Linking allows modularization

- at compile time
- at load-time
- at run-time

■ Linking involves two main operations

- symbol resolution: map symbols to a unique memory address across the entire program
- symbol relocation: ensure symbols refer to the designated memory address

■ Executable and linkable format

- common binary format of executables, libraries, and compiled object files
- contains all necessary information to support static and dynamic linking and loading
- very versatile and used on (almost) all platforms

Program Execution

■ Symbol Resolution

- first step of linking & loading
- for each use of a symbol, determine what defined symbol the use refers to
- (used to be) a source of subtle errors
- global, local, and external, strong and weak symbols
- symbol data to ELF section assignment: .text, .data, .bss

■ Symbol Relocation

- merge object files into a single executable/library
- relocate at-compile time unknown addresses/offsets to point to correct memory location
- two relocations of interest to us
 - ▶ PC-relative ($R_X86_64_PC32/PLT32$)
 - ▶ Absolute ($R_X86_64_32[S]/64$)

Program Execution

■ Libraries

- “packages” of functions commonly used together

■ Static Libraries

- “concatenation” of relocatable object files into an archive (hence *.a)
- at link time, “copy-paste” referenced object files into executable
- disadvantages: code size increase, library updates requires re-linking

■ Dynamic (Shared) Libraries

- library linked to executable at load/run-time
- allows sharing of code between different processes
- no recompilation necessary

■ Library Interpositioning

- intercept calls to system libraries
- a potential security risk